

STAT40800 Data Programming with Python - Final Project

Name & Student ID: Clementine Surya & 22200226

For this project, I analyse data from the Allen Institute, a bioscience research institute located in Seattle. Three datasets have been provided for this project:

- neurons group 1.csv: contains information about the morphology of a group of brain cell, known as neurons.
- neurons group 2.csv: contains information about the morphology of a second group of neurons.
- neurons additional measurements.csv: contains additional morphology measurements for the neurons included in neurons group 1.csv and neurons group 2.csv.

```
In [1]: #First I load in necessary packages for this project:
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy.random as npr
import scipy.stats as stats
import statsmodels.api as sm
import warnings
warnings.filterwarnings('ignore')
```

Question 1

(a) Load the neurons group 1.csv dataset into Python as a pandas DataFrame.

```
In [2]: #First of all, I load the neuron_group_1.csv dataset as
#a pandas DataFrame using below code
path = 'finalproject/'
df_neurons1 = pd.read_csv(path+'neurons_group_1.csv')

#Then, I check the first five row of the dataset
print(df_neurons1.head())
```

	id	average_diameter	overall_depth	overall_height	overall_width	\
0	484775243	0.195628	90.3529	548.798070	257.109717	
1	485996843	0.457635	87.0383	717.408343	199.214267	
2	486041253	0.295455	75.3286	584.083922	386.076695	
3	491119181	0.414033	89.0718	284.641670	239.492610	
4	491119245	0.201323	44.5237	302.038542	323.493562	

	soma_surface	total_length	total_surface	total_volume
0	128.269219	3658.629571	2252.681880	115.626135
1	430.635072	4158.819949	5944.196007	730.014704
2	502.033948	2667.618389	2472.795020	197.063796
3	383.828302	1543.941010	2008.838025	237.466517
4	120.229052	1621.871325	1027.220686	54.521240

(b) Inspect the data. How many neurons are included in this dataset? How many different measurements are included? Does this dataset contain any missing values?

There are 311 neurons entries with 8 measurements in neurons_group_1 dataset which are: 1) Average diameter, 2) Overall depth, 3) Overall height, 4) Overall width, 5) Soma surface, 6) Total length, 7) Total surface, and 8) Total volume.

There are 9 columns in this dataset and this dataset do not contain any missing values.

```
In [3]: #Below is the code that I use to show the information of the dataset
print(df_neurons1.info())
print(df_neurons1.shape)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 311 entries, 0 to 310
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    311 non-null    int64
1   average_diameter      311 non-null    float64
2   overall_depth         311 non-null    float64
3   overall_height        311 non-null    float64
4   overall_width         311 non-null    float64
5   soma_surface          311 non-null    float64
```

```

6   total_length      311 non-null    float64
7   total_surface     311 non-null    float64
8   total_volume      311 non-null    float64
dtypes: float64(8), int64(1)
memory usage: 22.0 KB
None
(311, 9)

```

```

In [4]: #Here, I count the missing values in each column of the dataset and
#as we can see there is no missing values in this data set
df_neurons1.isna().sum()

```

```

Out[4]: id                0
average_diameter         0
overall_depth            0
overall_height           0
overall_width            0
soma_surface             0
total_length             0
total_surface            0
total_volume             0
dtype: int64

```

(c) Perform an exploratory data analysis, creating both numerical and graphical summaries of the data. Discuss and interpret your results.

```

In [5]: #Here, I show the numerical summary for each of the neuron measurements
print(df_neurons1.describe())

```

	id	average_diameter	overall_depth	overall_height	\
count	3.110000e+02	311.000000	311.000000	311.000000	
mean	5.885866e+08	0.421175	91.967024	523.516774	
std	8.490396e+07	0.159435	35.658760	299.256684	
min	4.847752e+08	0.053899	22.680000	82.836871	
25%	4.961239e+08	0.322128	64.574450	328.964661	
50%	5.912744e+08	0.415613	86.212200	469.151885	
75%	6.568502e+08	0.527572	115.566350	651.885452	
max	8.460831e+08	1.156730	183.960000	1928.118350	

	overall_width	soma_surface	total_length	total_surface	total_volume
count	311.000000	311.000000	311.000000	311.000000	311.000000
mean	320.548089	361.849689	3792.940198	5492.737135	808.488483
std	124.039953	253.912293	2775.993259	5719.713426	1190.466808
min	49.173247	2.895610	251.987893	402.242787	4.309842
25%	233.480547	176.194233	2237.031837	2253.626506	197.071679
50%	291.532148	311.568275	2946.858866	3636.755377	407.180171
75%	392.534334	478.796933	4085.692910	5895.663836	833.681728
max	827.752239	1283.720986	15697.415190	37182.284100	8482.061401

Then, I make interpretation based on the numerical summary I obtain:

- The average diameter has a mean of 0.42 and standard deviation of 0.16. The median (0.42) is approximately equal to the mean, indicating that its distribution is symmetric.
- The overall depth has a mean of 91.97 and standard deviation of 35.66. The median (86.21) is lower than the mean, indicating that the distribution is right skewed.
- The overall height has a mean of 523.52 and standard deviation of 299.26. The median (469.15) is lower than the mean, indicating that the distribution is right skewed.
- The overall width has a mean of 320.55 and standard deviation of 124.04. The median (291.53) is lower than the mean, indicating that the distribution is right skewed.
- The soma surface has mean of 361.85 and standard deviation of 253.91. The median (311.57) is lower than the mean, indicating that the distribution is right skewed.
- The total length has mean of 3792.94 and standard deviation of 2775.99. The median (2946.86) is lower than the mean, indicating that the distribution is right skewed.
- The total surface has mean of 5492.74 and standard deviation of 5719.71. The median (3636.76) is significantly lower than the mean, indicating that the distribution is right skewed.
- The total volume has mean of 808.49 and standard deviation of 1190.47. The median (407.18) is significantly lower than the mean, indicating that the distribution is right skewed.

```

In [6]: #Below, I show the graphical summaries of this dataset measurements
sns.set(style="whitegrid")

fig, axes = plt.subplots(2, 4, figsize = (18, 12))

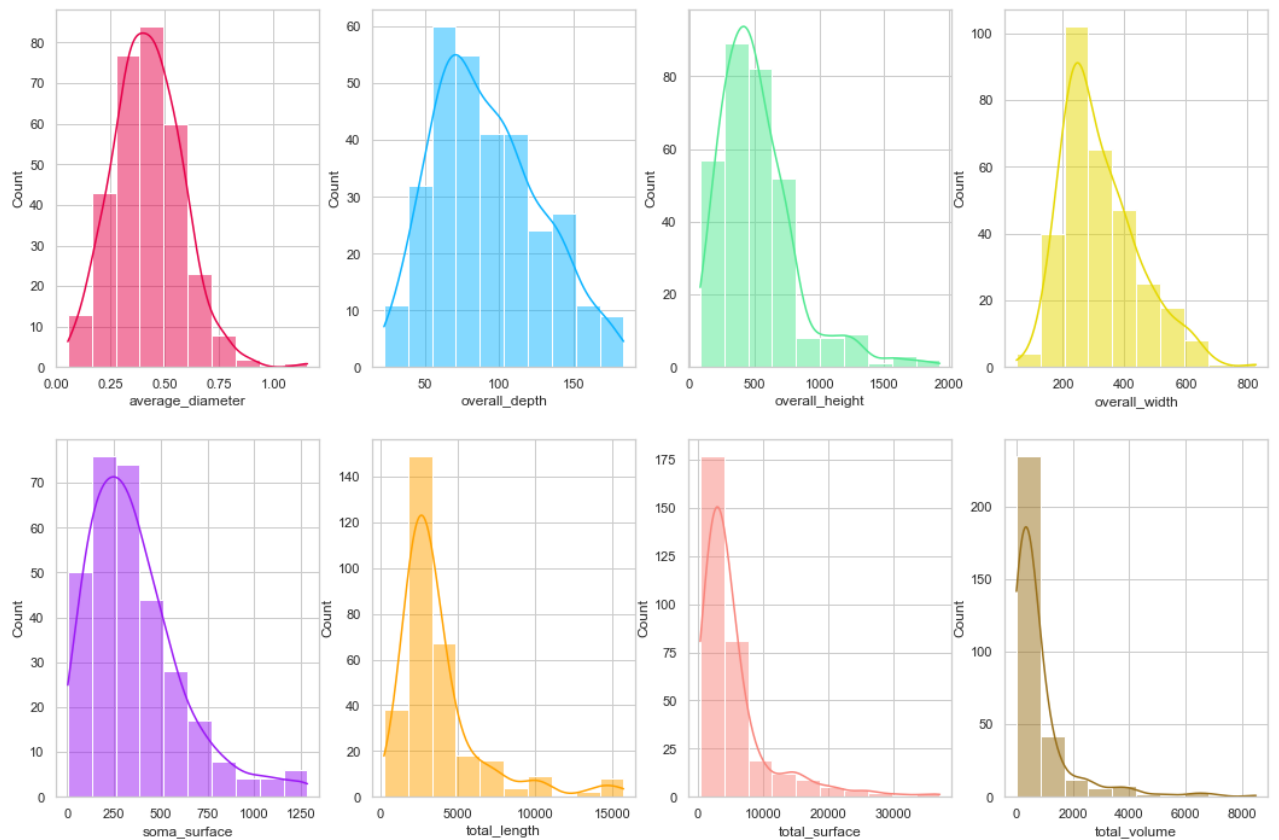
```

```

sns.histplot(data=df_neurons1,
             x="average_diameter",
             kde=True,
             color="#e60049",
             ax=axes[0, 0],
             bins =10).set(xlabel = "average_diameter")
sns.histplot(data=df_neurons1,
             x="overall_depth",
             kde=True,
             color="#0bb4ff",
             ax=axes[0, 1],
             bins =10).set(xlabel = "overall_depth")
sns.histplot(data=df_neurons1,
             x="overall_height",
             kde=True,
             color="#50e991",
             ax=axes[0, 2],
             bins =10).set(xlabel = "overall_height")
sns.histplot(data=df_neurons1,
             x="overall_width",
             kde=True,
             color="#e6d800",
             ax=axes[0, 3],
             bins =10).set(xlabel = "overall_width")
sns.histplot(data=df_neurons1,
             x="soma_surface",
             kde=True,
             color="#9b19f5",
             ax=axes[1, 0],
             bins =10).set(xlabel = "soma_surface")
sns.histplot(data=df_neurons1,
             x="total_length",
             kde=True,
             color="#ffa300",
             ax=axes[1, 1],
             bins =10).set(xlabel = "total_length")
sns.histplot(data=df_neurons1,
             x="total_surface",
             kde=True,
             color="#f9857c",
             ax=axes[1, 2],
             bins =10).set(xlabel = "total_surface")
sns.histplot(data=df_neurons1,
             x="total_volume",
             kde=True,
             color="#99701b",
             ax=axes[1, 3],
             bins =10).set(xlabel = "total_volume")

plt.show()

```



In [7]:

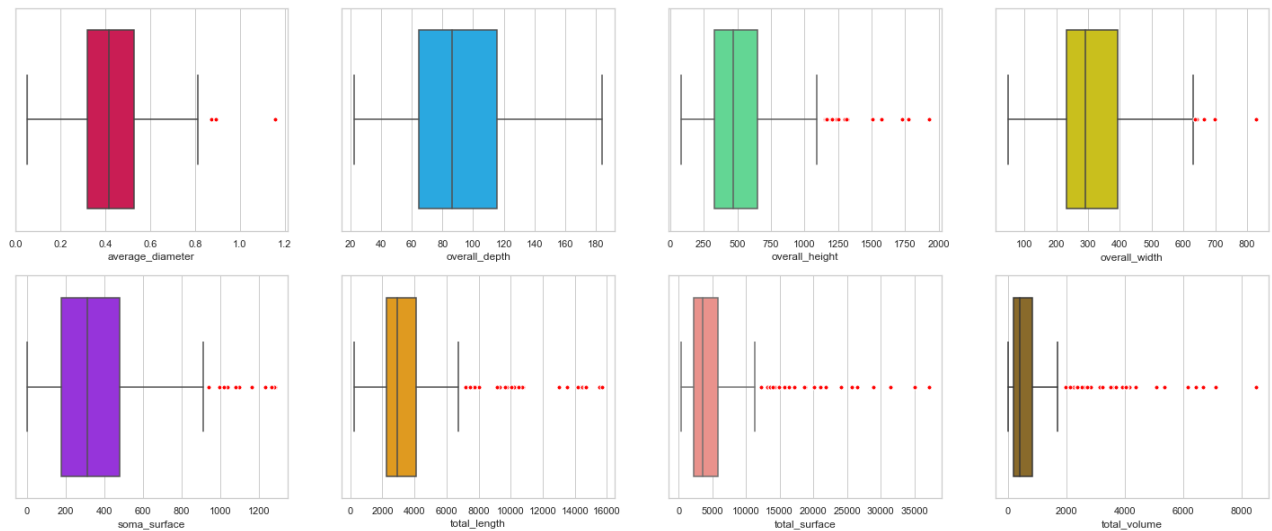
```
#I create boxplot neuron measurements summary with below code:
sns.set(style="whitegrid")

#Change the outlier marker to color red
red_circle = dict(markerfacecolor = 'red',
                  marker='o',
                  markeredgcolor='white')

fig,axs = plt.subplots(2,4, figsize = (25,10))
sns.boxplot(data=df_neurons1,
            x="average_diameter",
            color="#e60049",
            flierprops = red_circle,
            ax=axs[0, 0]).set(xlabel = "average_diameter")
sns.boxplot(data=df_neurons1,
            x="overall_depth",
            color="#0bb4ff",
            flierprops = red_circle,
            ax=axs[0, 1]).set(xlabel = "overall_depth")
sns.boxplot(data=df_neurons1,
            x="overall_height",
            color="#50e991",
            flierprops = red_circle,
            ax=axs[0, 2]).set(xlabel = "overall_height")
sns.boxplot(data=df_neurons1,
            x="overall_width",
            color="#e6d800",
            flierprops = red_circle,
            ax=axs[0, 3]).set(xlabel = "overall_width")
sns.boxplot(data=df_neurons1,
            x="soma_surface",
            color="#9b19f5",
            flierprops = red_circle,
            ax=axs[1, 0]).set(xlabel = "soma_surface")
sns.boxplot(data=df_neurons1,
            x="total_length",
            color="#ffa300",
            flierprops = red_circle,
            ax=axs[1, 1]).set(xlabel = "total_length")
sns.boxplot(data=df_neurons1,
            x="total_surface",
            color="#f9857c",
            flierprops = red_circle,
            ax=axs[1, 2]).set(xlabel = "total_surface")
sns.boxplot(data=df_neurons1,
            x="total_volume",
```

```
color="#99701b",
flierprops = red_circle,
ax=axes[1, 3]).set(xlabel = "total_volume")

plt.show()
```



Here, I do interpretation based on the graphical summary for each of the neurons measurement and connect it with the numerical summary. To sum up:

- Average diameter measurement has a symmetrical distribution, which is also shown in the numerical summary part, where the mean and median are both around 0.42. There are several outliers located above the maximum value in this distribution. The spread of the middle distribution is between 0.32 and 0.53 (IQR range of 0.21) and it is ranging from 0.05 to 1.16 (range value of 1.11).
- Overall depth measurement has a right skewed distribution, which is also shown in the numerical summary part, where the mean of overall depth measurement (91.97) is bigger than its median (86.21). The spread of the middle distribution is between 64.57 and 115.57 (IQR range of 51) and it is ranging from 22.68 to 183.96 (range value of 161.28).
- Overall height measurement has a right skewed distribution, which is also shown in the numerical summary part, where the mean of overall height measurement (523.52) is bigger than its median (469.15). There are several outliers located above the maximum value in this distribution. The spread of the middle distribution is between 328.96 and 651.89 (IQR range of 322.93) and it is ranging from 82.84 to 1928.12 (range value of 1845.28).
- Overall width measurement has a right skewed distribution, which is also shown in the numerical summary part, where the mean of overall width measurement (320.55) is bigger than its median (291.53). There are several outliers located above the maximum value in this distribution. The spread of the middle distribution is between 233.48 and 392.53 (IQR range of 159.05) and it is ranging from 49.17 to 827.75 (range value of 778.58).
- Soma surface measurement has a right skewed distribution, which is also shown in the numerical summary part, where the mean of soma surface measurement (361.85) is bigger than its median (311.57). There are many outliers located above the maximum value in this distribution. The spread of the middle distribution is between 176.19 and 478.80 (IQR range of 302.61) and it is ranging from 2.90 to 1283.72 (range value of 1280.82).
- Total length measurement has a right skewed distribution, which is also shown in the numerical summary part, where the mean of total length measurement (3792.94) is bigger than its median (2946.86). There are many outliers located above the maximum value in this distribution. The spread of the middle distribution is between 2237.03 and 4085.69 (IQR range of 1848.66) and it is ranging from 251.99 and 15697.42 (range value of 15445.43).
- Total surface has a right skewed distribution, which is also shown in the numerical summary part, where the mean of total surface measurement (5492.74) is bigger than its median (3636.76). There are many outliers located above the maximum value in this distribution. The spread of the middle distribution is between 2253.63 and 5895.66 (IQR range of 3642.03) and it is ranging from 402.24 and 37182.28 (range value 36780.04).
- Total volume has a right skewed distribution, which is also shown in the numerical summary part, where the mean of total volume measurement (808.49) is bigger than its median (407.18). There are many outliers located above the maximum value in this distribution. The spread of the middle distribution is between 197.07 and 833.68 (IQR range of 636.61) and it is ranging from 4.31 and 8482.06 (range value 8477.75).

Question 2

(a) Load the neurons group 2.csv dataset into Python as a pandas DataFrame.

```
In [8]: #I load the neuron_group_2.csv dataset as
#a pandas DataFrame using below code
path = 'finalproject/'
df_neurons2 = pd.read_csv(path+'neurons_group_2.csv')

#I show the first 5 rows of the dataset
print(df_neurons2.head())
```

	id	average_diameter	overall_depth	overall_height	overall_width	\
0	397905347	0.316091	117.5429	585.602322	287.122628	
1	491119234	0.331268	81.9012	461.280515	275.146120	
2	491119269	0.139015	57.5697	324.422347	280.851229	
3	491119394	0.230412	76.0357	368.298267	251.377567	
4	491119419	0.321163	98.8344	417.890620	193.590563	

	soma_surface	total_length	total_surface	total_volume
0	268.777679	3498.090031	3523.606841	306.290931
1	551.788645	2008.302439	2097.688550	188.431435
2	50.092109	1774.258366	776.076427	25.733000
3	244.457685	1650.188964	1198.531518	75.823080
4	252.423672	2066.369729	2085.789861	180.446414

(b) Inspect the data. How many neurons are included in this dataset? Are the measurements the same as those in neurons group 1.csv?

There are 390 neurons entries and 8 measurements in neuron_group_2 dataset. The measurements are the same as those in neurons group1.csv, where there are 8 measurements: 1) Average diameter, 2) Overall depth, 3) Overall height, 4) Overall width, 5) Soma surface, 6) Total length, 7) Total surface, and 8) Total volume.

There are 9 columns in this dataset and this dataset do not contain any missing values.

```
In [9]: #Below is the code that I use to show the information of the dataset
print(df_neurons2.info())
print(df_neurons2.shape)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 390 entries, 0 to 389
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     390 non-null    int64
1   average_diameter      390 non-null    float64
2   overall_depth         390 non-null    float64
3   overall_height        390 non-null    float64
4   overall_width         390 non-null    float64
5   soma_surface          390 non-null    float64
6   total_length          390 non-null    float64
7   total_surface         390 non-null    float64
8   total_volume          390 non-null    float64
dtypes: float64(8), int64(1)
memory usage: 27.5 KB
None
(390, 9)
```

```
In [10]: #Here, I count the missing values in each column of the dataset and
#as we can see there is no missing values in this data set
df_neurons2.isna().sum()
```

```
Out[10]: id                     0
average_diameter      0
overall_depth         0
overall_height        0
overall_width         0
soma_surface          0
total_length          0
total_surface         0
total_volume          0
dtype: int64
```

(c) Perform a t-test, for each of the measurements, to test whether any of the neuron properties differ between the group 1 and group 2. Use a significance level of $\alpha = 0.01$. Display the t-score and p-value for each measurement. Clearly state the conclusion of your tests and explain your reasoning.

In this section, I perform t-test for each measurements. A t-test is an inferential statistic used to determine if there is a significant difference between the means of two groups and how they are related¹. Hence, I would like to test whether any of the neuron properties differ between the group 1 and group 2. I use significance level of $\alpha = 0.01$.

The hypothesis testing that we use is:

H_0 : There is no difference in the neuron measurement means between the two groups

H_A : There is a difference in the neuron means between the two groups

We can see from the T-test result below that all the neurons measurement have p-value greater than $\alpha = 0.01$, we fail to reject the null hypothesis. Thus, we can conclude that there is no difference in the measurements means between neurons1 and neurons2 dataset.

T-test for average diameter measurement

H_0 : There is no difference in the average diameter means between the two groups

H_A : There is a difference in the average diameter means between the two groups

As we can see from the result below, the T-score is 1.197 and p-value= 0.23 which is greater than $\alpha = 0.01$. Thus, we fail to reject the null hypothesis. This also indicates that there is no difference in the average diameter means between neurons1 and neurons2 dataset.

```
In [11]: stats.ttest_ind(df_neurons1['average_diameter'], df_neurons2['average_diameter'])
```

```
Out[11]: Ttest_indResult(statistic=-1.1970075781748382, pvalue=0.23170939849935634)
```

T-test for overall depth measurement

H_0 : There is no difference in the overall depth means between the two groups

H_A : There is a difference in the overall depth means between the two groups

As we can see from the result below, the T-score is 1.19 and p-value= 0.24 which is greater than $\alpha = 0.01$. Thus, we fail to reject the null hypothesis. This also indicates that there is no difference in the overall depth means between neurons1 and neurons2 dataset.

```
In [12]: stats.ttest_ind(df_neurons1['overall_depth'], df_neurons2['overall_depth'])
```

```
Out[12]: Ttest_indResult(statistic=-1.1857230913955603, pvalue=0.23613458215666971)
```

T-test for overall height measurement

H_0 : There is no difference in the overall height means between the two groups

H_A : There is a difference in the overall height means between the two groups

As we can see from the result below, the T-score is 0.59 and p-value= 0.56 which is greater than $\alpha = 0.01$. Thus, we fail to reject the null hypothesis. This also indicates that there is no difference in the overall height means between neurons1 and neurons2 dataset.

```
In [13]: stats.ttest_ind(df_neurons1['overall_height'], df_neurons2['overall_height'])
```

```
Out[13]: Ttest_indResult(statistic=0.5860348440044614, pvalue=0.5580412742530052)
```

T-test for overall width measurement

H_0 : There is no difference in the overall width means between the two groups

H_A : There is a difference in the overall width means between the two groups

As we can see from the result below, the T-score is 0.92 and p-value= 0.36 which is greater than $\alpha = 0.01$. Thus, we fail to reject the null hypothesis. This also indicates that there is no difference in the overall width measurement between neurons1 and neurons2 dataset.

```
In [14]: stats.ttest_ind(df_neurons1['overall_width'], df_neurons2['overall_width'])
```

```
Out[14]: Ttest_indResult(statistic=-0.9231920391791176, pvalue=0.3562256919514518)
```

T-test for soma surface measurement

H_0 : There is no difference in the soma surface means between the two groups

H_A : There is a difference in the soma surface means between the two groups

As we can see from the result below, the T-score is 0.55 and p-value= 0.58 which is greater than $\alpha = 0.01$. Thus, we fail to reject the null hypothesis. This also indicates that there is no difference in the soma surface means between neurons1 and neurons2 dataset.

```
In [15]: stats.ttest_ind(df_neurons1['soma_surface'], df_neurons2['soma_surface'])
```

```
Out[15]: Ttest_indResult(statistic=-0.547797643314149, pvalue=0.5840057333717408)
```

T-test for total length measurement

H_0 : There is no difference in the total length means between the two groups

H_A : There is a difference in the total length means between the two groups

As we can see from the result below, the T-score is 0.33 and p-value= 0.74 which is greater than $\alpha = 0.01$. Thus, we fail to reject the null hypothesis. This also indicates there is no difference in the total length means between neurons1 and neurons2 dataset.

```
In [16]: stats.ttest_ind(df_neurons1['total_length'], df_neurons2['total_length'])
```

```
Out[16]: Ttest_indResult(statistic=-0.332599650604377, pvalue=0.7395362277560685)
```

T-test for total surface measurement

H_0 : There is no difference in the total surface means between the two groups

H_A : There is a difference in the total surface means between the two groups

As we can see from the result below, the T-score is 0.82 and p-value= 0.42 which is greater than $\alpha = 0.01$. Thus, we fail to reject the null hypothesis. This also indicates there is no difference in the total surface means between neurons1 and neurons2 dataset.

```
In [17]: stats.ttest_ind(df_neurons1['total_surface'], df_neurons2['total_surface'])
```

```
Out[17]: Ttest_indResult(statistic=-0.8152397825540432, pvalue=0.4152127535605644)
```

T-test for total volume measurement

H_0 : There is no difference in the total volume means between the two groups

H_A : There is a difference in the total volume means between the two groups

As we can see from the result below, the T-score is 0.93 and p-value is 0.35 which is greater than $\alpha = 0.01$. Thus, we fail to reject the null hypothesis. This also indicates there is no difference in the total volume means between neurons1 and neurons2 dataset.

```
In [18]: stats.ttest_ind(df_neurons1['total_volume'], df_neurons2['total_volume'])
```

```
Out[18]: Ttest_indResult(statistic=-0.9313735524535135, pvalue=0.35198192808008355)
```

Question 3

(a) Load the neurons additional measurements.csv into Python and combine all three datasets into a single DataFrame.

After loading the neurons additional measurement dataset with below code, we can see that there are 694 neurons and 13 columns in this dataset with 12 neuron measurements. All the measurements are different from the measurements in neurons1 and neurons2 datasets. The 12 measurements in this additional dataset are average bifurcation angle local, average contraction, average fragmentation, average parent daughter ratio, max branch order, max euclidean distance, max path distance, number bifurcations, number branches, number nodes, number stems, and number tips.

```
In [19]: #I load the neurons_additional_measurements.csv dataset as a pandas DataFrame using below code
```



```
path = 'finalproject/'
df_neurons_additional = pd.read_csv(path+'neurons_additional_measurements.csv')

#I show the 5 five rows of this dataset
print(df_neurons_additional.head())
```

```
      id  average_bifurcation_angle_local  average_contraction \
0  491119743                82.727781          0.864267
1  546781359                82.506680          0.903890
2  537042261                77.536678          0.863104
3  689123605                76.583222          0.900537
4  657879305                72.019250          0.873518

      average_fragmentation  average_parent_daughter_ratio  max_branch_order \
0                20.723077                0.964510                6
1               105.277778                0.862183                3
2                73.666667                0.926633                6
3                95.979167                0.942049               11
4                47.535714                1.000000                5

      max_euclidean_distance  max_path_distance  number_bifurcations \
0                99.779724        126.593790                33
1               432.383110        496.831994                9
2               373.630444        436.958952               21
3               943.382549        989.448318               24
4               186.218009        221.639502               14

      number_branches  number_nodes  number_stems  number_tips
0                73          1470            7          40
1                23          2011            5          14
2                46          3137            4          25
3                52          4652            4          28
4                32          1406            4          18
```

In [20]:

```
#Show the information of this additional dataset and
#the shape of the dataset using below code
print(df_neurons_additional.info())
print(df_neurons_additional.shape)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 694 entries, 0 to 693
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    694 non-null   int64
1   average_bifurcation_angle_local      694 non-null   float64
2   average_contraction                  694 non-null   float64
3   average_fragmentation                 694 non-null   float64
4   average_parent_daughter_ratio        694 non-null   float64
5   max_branch_order                     694 non-null   int64
6   max_euclidean_distance                694 non-null   float64
7   max_path_distance                    694 non-null   float64
8   number_bifurcations                  694 non-null   int64
9   number_branches                      694 non-null   int64
10  number_nodes                         694 non-null   int64
11  number_stems                         694 non-null   int64
12  number_tips                          694 non-null   int64
dtypes: float64(6), int64(7)
memory usage: 70.6 KB
None
(694, 13)
```

Here, I combine neurons1 and neurons2 dataset together. We can see that there are 701 neurons entries and 9 columns in the combined dataset of neurons1 and neurons2 with 8 neurons measurement that we have already seen in part 1 and part 2 questions.

In [21]:

```
#I use append function to combine neuron_group_1
#and neuron_group_2 datasets
df_neurons1_neurons2 = df_neurons1.append(df_neurons2)

#Here, I show the information and dimension of the combine dataset
#of neuron_group_1 and neuron_group_2 datasets
print(df_neurons1_neurons2.info())
print(df_neurons1_neurons2.shape)
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 701 entries, 0 to 389
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	id	701 non-null	int64
1	average_diameter	701 non-null	float64
2	overall_depth	701 non-null	float64
3	overall_height	701 non-null	float64
4	overall_width	701 non-null	float64
5	soma_surface	701 non-null	float64
6	total_length	701 non-null	float64
7	total_surface	701 non-null	float64
8	total_volume	701 non-null	float64

dtypes: float64(8), int64(1)

memory usage: 54.8 KB

None

(701, 9)

After that, I merge the df_neurons1_neurons2 dataset with df_neurons_additional dataset and make a new dataset that contains all of the data called df_neurons_all.

```
In [22]: #Merge df_neurons1_neurons2 dataset with df_neurons_additional dataset using "inner"
#This keeps only the common values in both
#df_neurons1_neurons2 and df_neurons_additional dataset
#for the merged data
df_neurons_all = pd.merge(df_neurons1_neurons2,
                          df_neurons_additional,on=['id'],how = 'inner')

#Show the 5 row of this dataset
print(df_neurons_all.head())
```

	id	average_diameter	overall_depth	overall_height	overall_width	\
0	484775243	0.195628	90.3529	548.798070	257.109717	
1	485996843	0.457635	87.0383	717.408343	199.214267	
2	486041253	0.295455	75.3286	584.083922	386.076695	
3	491119181	0.414033	89.0718	284.641670	239.492610	
4	491119245	0.201323	44.5237	302.038542	323.493562	

	soma_surface	total_length	total_surface	total_volume	\
0	128.269219	3658.629571	2252.681880	115.626135	
1	430.635072	4158.819949	5944.196007	730.014704	
2	502.033948	2667.618389	2472.795020	197.063796	
3	383.828302	1543.941010	2008.838025	237.466517	
4	120.229052	1621.871325	1027.220686	54.521240	

	average_bifurcation_angle_local	...	average_fragmentation	\
0	78.026948	...	48.081967	
1	86.349668	...	52.416667	
2	71.911353	...	73.821429	
3	72.177287	...	42.266667	
4	70.758367	...	39.733333	

	average_parent_daughter_ratio	max_branch_order	max_euclidean_distance	\
0	0.866582	7	476.921223	
1	0.950502	15	650.637431	
2	0.816243	5	477.440600	
3	0.822614	4	212.448082	
4	0.803137	3	242.351930	

	max_path_distance	number_bifurcations	number_branches	number_nodes	\
0	550.691106	31	70	3041	
1	957.710930	30	67	3271	
2	569.974920	14	32	2151	
3	249.521471	16	35	1292	
4	258.868561	15	36	1397	

	number_stems	number_tips
0	8	39
1	7	37
2	4	18
3	3	19
4	6	21

[5 rows x 21 columns]

(b) Comment on the dimensions of the combined dataset. Are all of the neurons from group 1 and 2 included in the dataset neurons_additional measurements.csv?

The dimensions of the combined dataset are 694 neurons entries and 21 column. From above part, we know that the neurons entries from neuron_group_1 are 311 and the neurons entries from neuron_group_2 are 390. The total neuron entries of neuron_group_1 and neuron_group_2 are 701. This means that there are 7 neurons from group 1 and 2 that are

not included in the dataset of neurons additional measurements.csv. With the result shown below, we know that there are 5 neurons and 2 neurons that are not included in the combined dataset from neuron_group_1 and neuron_group_2 respectively.

We need to take note that this combined dataset does not contain any missing values because of the "inner" function when merging the dataset in previous section. Hence, this keeps only the common neurons "id" values in the combined dataset (proven below that there is no missing value).

In [23]:

```
#Show the information and dimension of the dataset
print(df_neurons_all.info())
print(df_neurons_all.shape)
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 694 entries, 0 to 693
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    694 non-null    int64
1   average_diameter                     694 non-null    float64
2   overall_depth                        694 non-null    float64
3   overall_height                       694 non-null    float64
4   overall_width                        694 non-null    float64
5   soma_surface                         694 non-null    float64
6   total_length                         694 non-null    float64
7   total_surface                        694 non-null    float64
8   total_volume                         694 non-null    float64
9   average_bifurcation_angle_local      694 non-null    float64
10  average_contraction                  694 non-null    float64
11  average_fragmentation                694 non-null    float64
12  average_parent_daughter_ratio        694 non-null    float64
13  max_branch_order                     694 non-null    int64
14  max_euclidean_distance                694 non-null    float64
15  max_path_distance                    694 non-null    float64
16  number_bifurcations                  694 non-null    int64
17  number_branches                      694 non-null    int64
18  number_nodes                         694 non-null    int64
19  number_stems                         694 non-null    int64
20  number_tips                          694 non-null    int64
dtypes: float64(14), int64(7)
memory usage: 119.3 KB
None
(694, 21)
```

In [24]:

```
#Check any neurons in dataset 1 and 2
#that are not included in additional dataset
check_group1 = df_neurons1['id'].isin(df_neurons_additional['id'])
missing_group1 = check_group1.value_counts()[False]

check_group2 = df_neurons2['id'].isin(df_neurons_additional['id'])
missing_group2 = check_group2.value_counts()[False]

print("Number of neurons not included from Group 1:", missing_group1)
print("Number of neurons not included from Group 2:", missing_group2)
```

```
Number of neurons not included from Group 1: 5
Number of neurons not included from Group 2: 2
```

In [25]:

```
df_neurons_all.isna().sum()
```

Out[25]:

```
id                                0
average_diameter                  0
overall_depth                     0
overall_height                    0
overall_width                     0
soma_surface                      0
total_length                      0
total_surface                     0
total_volume                      0
average_bifurcation_angle_local   0
average_contraction               0
average_fragmentation             0
average_parent_daughter_ratio     0
max_branch_order                  0
max_euclidean_distance            0
max_path_distance                 0
number_bifurcations               0
```

```

number_branches      0
number_nodes         0
number_stems         0
number_tips          0
dtype: int64

```

(c) Compute the Pearson correlation coefficient between each of the measurements and identify which morphological features are strongly correlated. List the four most strongly correlated pairs.

In this part, I compute the Pearson correlation coefficient between each of the measurements and identify which morphological features are strongly correlated.

The result below shown that, the four most strongly correlated pairs are number nodes and total length (correlation: 0.998333), number tips and number branches (0.998230), number branches and number bifuractions (0.997993), and number tips and number bifuractions (0.992465).

In [26]:

```

neurons_numeric = df_neurons_all.drop(['id'],axis=1)
neurons_cor = neurons_numeric.corr()
print(neurons_cor)

```

	average_diameter	overall_depth	\
average_diameter	1.000000	0.229786	
overall_depth	0.229786	1.000000	
overall_height	0.327179	0.527627	
overall_width	0.367010	0.515622	
soma_surface	0.589195	0.238661	
total_length	0.352461	0.566376	
total_surface	0.619391	0.482570	
total_volume	0.658883	0.405527	
average_bifurcation_angle_local	0.039447	-0.067234	
average_contraction	-0.262180	-0.168192	
average_fragmentation	0.353716	0.420348	
average_parent_daughter_ratio	0.274922	0.004063	
max_branch_order	0.172831	0.496938	
max_euclidean_distance	0.300276	0.542190	
max_path_distance	0.313598	0.581248	
number_bifurcations	0.164862	0.434871	
number_branches	0.151537	0.434845	
number_nodes	0.345626	0.552837	
number_stems	-0.131749	0.166980	
number_tips	0.138467	0.433055	

	overall_height	overall_width	soma_surface	\
average_diameter	0.327179	0.367010	0.589195	
overall_depth	0.527627	0.515622	0.238661	
overall_height	1.000000	0.472310	0.306192	
overall_width	0.472310	1.000000	0.400587	
soma_surface	0.306192	0.400587	1.000000	
total_length	0.660042	0.680065	0.428615	
total_surface	0.618149	0.660372	0.549134	
total_volume	0.540829	0.579829	0.537021	
average_bifurcation_angle_local	-0.170636	-0.088734	-0.062920	
average_contraction	-0.011630	-0.070939	-0.049990	
average_fragmentation	0.546112	0.655019	0.330257	
average_parent_daughter_ratio	-0.057364	0.107348	0.109646	
max_branch_order	0.614627	0.269896	0.155667	
max_euclidean_distance	0.954598	0.482814	0.263454	
max_path_distance	0.941420	0.502016	0.265640	
number_bifurcations	0.432410	0.370920	0.266762	
number_branches	0.425560	0.365143	0.263538	
number_nodes	0.655673	0.673310	0.428377	
number_stems	0.066929	0.061535	0.055295	
number_tips	0.417486	0.358484	0.259445	

	total_length	total_surface	total_volume	\
average_diameter	0.352461	0.619391	0.658883	
overall_depth	0.566376	0.482570	0.405527	
overall_height	0.660042	0.618149	0.540829	
overall_width	0.680065	0.660372	0.579829	
soma_surface	0.428615	0.549134	0.537021	
total_length	1.000000	0.913416	0.790274	
total_surface	0.913416	1.000000	0.958704	
total_volume	0.790274	0.958704	1.000000	
average_bifurcation_angle_local	-0.143934	-0.124021	-0.113610	
average_contraction	-0.064172	-0.136860	-0.168131	
average_fragmentation	0.456874	0.474200	0.436098	
average_parent_daughter_ratio	0.102789	0.183065	0.181217	
max_branch_order	0.599113	0.495238	0.401908	

max_euclidean_distance	0.641754	0.597350	0.520168
max_path_distance	0.654762	0.608415	0.531298
number_bifurcations	0.817617	0.684764	0.561337
number_branches	0.811752	0.674198	0.550372
number_nodes	0.998333	0.909903	0.787010
number_stems	0.234539	0.113077	0.058827
number_tips	0.803354	0.661868	0.538113

	average_bifurcation_angle_local	\
average_diameter	0.039447	
overall_depth	-0.067234	
overall_height	-0.170636	
overall_width	-0.088734	
soma_surface	-0.062920	
total_length	-0.143934	
total_surface	-0.124021	
total_volume	-0.113610	
average_bifurcation_angle_local	1.000000	
average_contraction	-0.256651	
average_fragmentation	-0.033428	
average_parent_daughter_ratio	0.144248	
max_branch_order	-0.057797	
max_euclidean_distance	-0.160161	
max_path_distance	-0.138639	
number_bifurcations	-0.104305	
number_branches	-0.107519	
number_nodes	-0.139077	
number_stems	-0.088756	
number_tips	-0.110187	

	average_contraction	average_fragmentation	\
average_diameter	-0.262180	0.353716	
overall_depth	-0.168192	0.420348	
overall_height	-0.011630	0.546112	
overall_width	-0.070939	0.655019	
soma_surface	-0.049990	0.330257	
total_length	-0.064172	0.456874	
total_surface	-0.136860	0.474200	
total_volume	-0.168131	0.436098	
average_bifurcation_angle_local	-0.256651	-0.033428	
average_contraction	1.000000	-0.232868	
average_fragmentation	-0.232868	1.000000	
average_parent_daughter_ratio	-0.196380	0.133164	
max_branch_order	-0.055551	0.041575	
max_euclidean_distance	-0.038668	0.463791	
max_path_distance	-0.095109	0.469273	
number_bifurcations	0.054079	-0.030394	
number_branches	0.059078	-0.037248	
number_nodes	-0.067872	0.457285	
number_stems	0.093060	-0.111443	
number_tips	0.063447	-0.043461	

	average_parent_daughter_ratio	\
average_diameter	0.274922	
overall_depth	0.004063	
overall_height	-0.057364	
overall_width	0.107348	
soma_surface	0.109646	
total_length	0.102789	
total_surface	0.183065	
total_volume	0.181217	
average_bifurcation_angle_local	0.144248	
average_contraction	-0.196380	
average_fragmentation	0.133164	
average_parent_daughter_ratio	1.000000	
max_branch_order	-0.136795	
max_euclidean_distance	-0.114534	
max_path_distance	-0.108647	
number_bifurcations	0.010898	
number_branches	-0.004622	
number_nodes	0.099655	
number_stems	-0.222174	
number_tips	-0.019046	

	max_branch_order	max_euclidean_distance	\
average_diameter	0.172831	0.300276	
overall_depth	0.496938	0.542190	
overall_height	0.614627	0.954598	
overall_width	0.269896	0.482814	
soma_surface	0.155667	0.263454	

total_length	0.599113	0.641754
total_surface	0.495238	0.597350
total_volume	0.401908	0.520168
average_bifurcation_angle_local	-0.057797	-0.160161
average_contraction	-0.055551	-0.038668
average_fragmentation	0.041575	0.463791
average_parent_daughter_ratio	-0.136795	-0.114534
max_branch_order	1.000000	0.678247
max_euclidean_distance	0.678247	1.000000
max_path_distance	0.706142	0.989598
number_bifurcations	0.712654	0.458671
number_branches	0.713865	0.456456
number_nodes	0.598367	0.638202
number_stems	0.295542	0.145440
number_tips	0.712331	0.452622

	max_path_distance	number_bifurcations \
average_diameter	0.313598	0.164862
overall_depth	0.581248	0.434871
overall_height	0.941420	0.432410
overall_width	0.502016	0.370920
soma_surface	0.265640	0.266762
total_length	0.654762	0.817617
total_surface	0.608415	0.684764
total_volume	0.531298	0.561337
average_bifurcation_angle_local	-0.138639	-0.104305
average_contraction	-0.095109	0.054079
average_fragmentation	0.469273	-0.030394
average_parent_daughter_ratio	-0.108647	0.010898
max_branch_order	0.706142	0.712654
max_euclidean_distance	0.989598	0.458671
max_path_distance	1.000000	0.477001
number_bifurcations	0.477001	1.000000
number_branches	0.474840	0.997993
number_nodes	0.650522	0.818667
number_stems	0.153536	0.361552
number_tips	0.470999	0.992465

	number_branches	number_nodes	number_stems \
average_diameter	0.151537	0.345626	-0.131749
overall_depth	0.434845	0.552837	0.166980
overall_height	0.425560	0.655673	0.066929
overall_width	0.365143	0.673310	0.061535
soma_surface	0.263538	0.428377	0.055295
total_length	0.811752	0.998333	0.234539
total_surface	0.674198	0.909903	0.113077
total_volume	0.550372	0.787010	0.058827
average_bifurcation_angle_local	-0.107519	-0.139077	-0.088756
average_contraction	0.059078	-0.067872	0.093060
average_fragmentation	-0.037248	0.457285	-0.111443
average_parent_daughter_ratio	-0.004622	0.099655	-0.222174
max_branch_order	0.713865	0.598367	0.295542
max_euclidean_distance	0.456456	0.638202	0.145440
max_path_distance	0.474840	0.650522	0.153536
number_bifurcations	0.997993	0.818667	0.361552
number_branches	1.000000	0.812985	0.419783
number_nodes	0.812985	1.000000	0.237684
number_stems	0.419783	0.237684	1.000000
number_tips	0.998230	0.804757	0.473023

	number_tips
average_diameter	0.138467
overall_depth	0.433055
overall_height	0.417486
overall_width	0.358484
soma_surface	0.259445
total_length	0.803354
total_surface	0.661868
total_volume	0.538113
average_bifurcation_angle_local	-0.110187
average_contraction	0.063447
average_fragmentation	-0.043461
average_parent_daughter_ratio	-0.019046
max_branch_order	0.712331
max_euclidean_distance	0.452622
max_path_distance	0.470999
number_bifurcations	0.992465
number_branches	0.998230
number_nodes	0.804757

```
number_stems          0.473023
number_tips           1.000000
```

In [27]:

```
#Retain upper triangular values of correlation matrix and
#make Lower triangular values Null
upper_corr_mat = neurons_cor.where(
    np.triu(np.ones(neurons_cor.shape), k=1).astype(bool))

#Convert to 1-D series and drop Null values
unique_corr_pairs = upper_corr_mat.unstack().dropna()

#Sort correlation pairs
sorted_mat = unique_corr_pairs.sort_values()
print(sorted_mat)
```

```
average_contraction      average_diameter      -0.262180
                        average_bifurcation_angle_local -0.256651
average_fragmentation    average_contraction    -0.232868
number_stems             average_parent_daughter_ratio -0.222174
average_parent_daughter_ratio average_contraction -0.196380
...
max_path_distance        max_euclidean_distance    0.989598
number_tips              number_bifurcations        0.992465
number_branches          number_bifurcations        0.997993
number_tips              number_branches           0.998230
number_nodes             total_length               0.998333
Length: 190, dtype: float64
```

(d) Creating scatterplot for the four most strongly correlated pairs.

We can see from the scatterplots below that number nodes & total length pair, number tips & number branches pair, number branches & number bifurcations pair, and number tips & number bifurcations pair are all having **strong positive correlation** between each other as expected.

In [28]:

```
sns.set(style="whitegrid")

fig,axs = plt.subplots(2,2, figsize = (18,12))

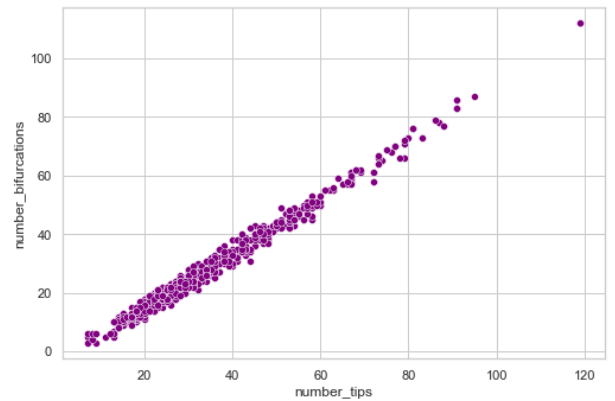
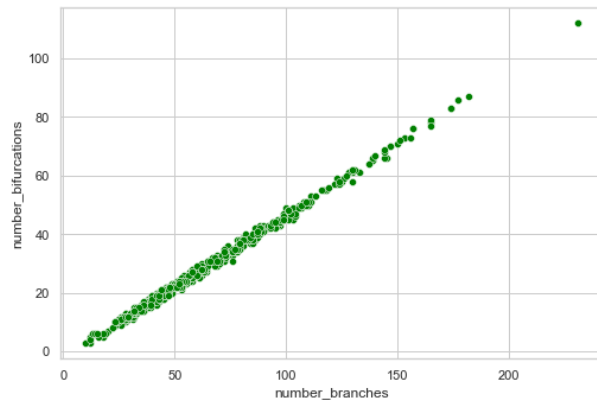
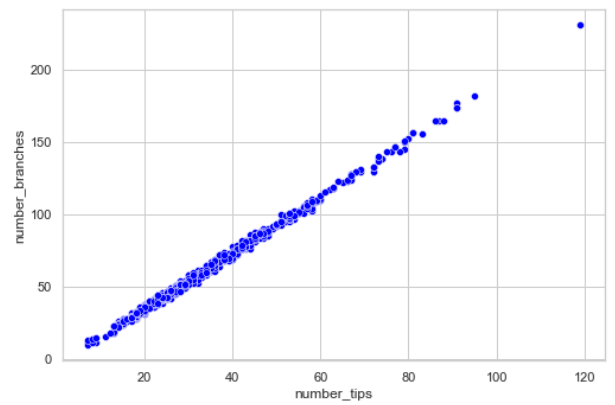
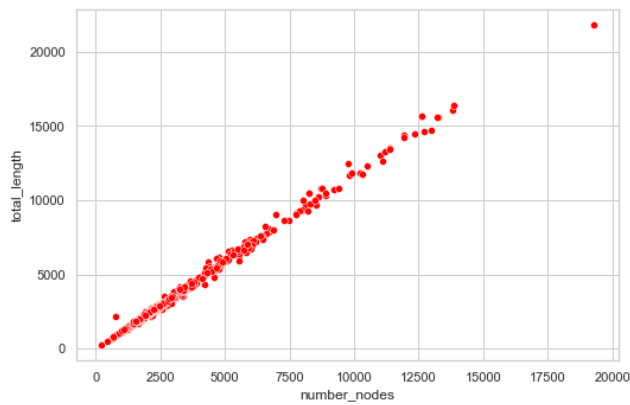
#Scatterplot between number nodes and total length measurements
sns.scatterplot(data = df_neurons_all,
    x = "number_nodes",
    y = "total_length",
    color = 'red',
    ax=axs[0, 0])

#Scatterplot between number tips and number branches measurements
sns.scatterplot(data=df_neurons_all,
    x="number_tips",
    y="number_branches",
    color = 'blue',
    ax=axs[0, 1])

#Scatterplot between number branches and number bifurcations measurements
sns.scatterplot(data=df_neurons_all,
    x="number_branches",
    y="number_bifurcations",
    color = 'green',
    ax=axs[1, 0])

#Scatterplot between number tips and number bifurcations measurements
sns.scatterplot(data=df_neurons_all,
    x="number_tips",
    y="number_bifurcations",
    color = 'purple',
    ax=axs[1, 1])

plt.show()
```

Question 4 - Linear regression to Predict the Total Surface Area of a Neuron

(a) Separate the data into response and predictor variables and standardise the predictor variables

In this section, I separate the data into response and predictor variables and standardise the predictor variables.

```
In [29]: #Create a Series y which contains the response variable (total surface)
y = df_neurons_all.total_surface

#Create a DataFrame X that excludes id and
#total surface and contains the explanatory variables:
#average diameter, overall depth, overall height, overall width,
#soma surface, total length, total volume
#average bifurcation angle local, average contraction,
#average fragmentation, average parent daughter ratio
#max branch order, max euclidean distance, max path distance,
#number bifurcations, number branches,
#number nodes, number stems, number tips
X = df_neurons_all.drop(['id', 'total_surface'], axis=1)

#Standardise the predictor variable
X_std = (X - X.mean()) / X.std()

#Add an intercept column to X in the first column
X_std.insert(0, 'intercept', 1)
X = X_std
```

(b) Fit a linear regression model and interpret the fitted model

Here I fit a linear regression model using statsmodels.api package that I have written in the "Load necessary package" part.

```
In [30]: mod = sm.OLS(y, X)
res = mod.fit()
print(res.summary())
```

```

=====
                    OLS Regression Results
=====
Dep. Variable:      total_surface      R-squared:                0.992
Model:              OLS               Adj. R-squared:           0.991
Method:             Least Squares      F-statistic:             4465.
Date:               Sun, 11 Dec 2022    Prob (F-statistic):       0.00
Time:               16:26:06           Log-Likelihood:          -5362.3

```

```

No. Observations:      694    AIC:                1.076e+04
Df Residuals:          675    BIC:                1.085e+04
Df Model:              18
Covariance Type:      nonrobust

```

	coef	std err	t	P> t	[0.025	0.975]
intercept	5705.7931	21.123	270.127	0.000	5664.319	5747.267
average_diameter	693.1540	36.058	19.223	0.000	622.354	763.953
overall_depth	-130.8024	31.030	-4.215	0.000	-191.730	-69.875
overall_height	-296.8748	90.839	-3.268	0.001	-475.235	-118.514
overall_width	25.8528	38.994	0.663	0.508	-50.712	102.418
soma_surface	56.9857	28.512	1.999	0.046	1.004	112.968
total_length	3956.4533	411.041	9.625	0.000	3149.380	4763.526
total_volume	2994.2462	48.152	62.183	0.000	2899.701	3088.792
average_bifurcation_angle_local	-5.6851	22.755	-0.250	0.803	-50.365	38.994
average_contraction	32.2640	26.690	1.209	0.227	-20.141	84.669
average_fragmentation	-95.7414	49.252	-1.944	0.052	-192.446	0.963
average_parent_daughter_ratio	54.8073	24.030	2.281	0.023	7.624	101.990
max_branch_order	-26.1737	42.612	-0.614	0.539	-109.841	57.494
max_euclidean_distance	636.0267	188.029	3.383	0.001	266.835	1005.218
max_path_distance	-346.6786	181.393	-1.911	0.056	-702.841	9.484
number_bifurcations	-3895.7825	6022.607	-0.647	0.518	-1.57e+04	7929.514
number_branches	1985.8946	3204.822	0.620	0.536	-4306.724	8278.513
number_nodes	-813.9945	408.079	-1.995	0.046	-1615.251	-12.738
number_stems	-458.2155	620.789	-0.738	0.461	-1677.125	760.694
number_tips	1897.7216	3070.338	0.618	0.537	-4130.840	7926.283
=====						
Omnibus:	611.210	Durbin-Watson:		1.913		
Prob(Omnibus):	0.000	Jarque-Bera (JB):		32896.241		
Skew:	-3.626	Prob(JB):		0.00		
Kurtosis:	35.940	Cond. No.		1.04e+16		
=====						

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 5.41e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

Interpreting the fitted models:

First of all, I interpret the coefficient:

The average diameter coefficient in the regression equation is 693.1540. This coefficient represents the mean increase of total surface for every additional unit in average diameter. If the average diameter increases by 1 unit, the average total surface increases by 693.1540 (holding all the other independent variables constant).

The overall depth coefficient in the regression equation is -130.8024. This coefficient represents the mean decrease of total surface for every additional unit in overall depth. If the overall depth increases by 1 unit, the average total surface decreases by 130.8024 (holding all the other independent variables constant).

The overall height coefficient in the regression equation is -296.8748. This coefficient represents the mean decrease of total surface for every additional unit in overall height. If the overall height increases by 1 unit, the average total surface decreases by 296.8748 (holding all the other independent variables constant).

The overall width coefficient in the regression equation is 25.8528. This coefficient represents the mean increase of total surface for every additional unit in overall width. If the overall width increases by 1 unit, the average total surface increases by 25.8528 (holding all the other independent variables constant).

The soma surface coefficient in the regression equation is 56.9857. This coefficient represents the mean increase of total surface for every additional unit in soma surface. If the soma surface increases by 1 unit, the average total surface increases by 56.9857 (holding all the other independent variables constant).

The total length coefficient in the regression equation is 3956.4533. This coefficient represents the mean increase of total surface for every additional unit in total length. If the total length increases by 1 unit, the average total surface increases by 3956.4533 (holding all the other independent variables constant).

The total volume coefficient in the regression equation is 2994.2462. This coefficient represents the mean increase of total surface for every additional unit in total volume. If the total volume increases by 1 unit, the average total surface increases by 2994.2462 (holding all the other independent variables constant).

The average bifurcation angle local coefficient in the regression equation is -5.6851. This coefficient represents the mean decrease of total surface for every additional unit in average bifurcation angle local. If the average bifurcation angle local by 1 unit, the average total surface decreases by 5.6851 (holding all the other independent variables constant).

The average contraction coefficient in the regression equation is 32.2640. This coefficient represents the mean increase of total surface for every additional unit in average contraction. If the average contraction increases by 1 unit, the average total surface increases by 32.2640 (holding all the other independent variables constant).

The average fragmentation coefficient in the regression equation is -95.7414. This coefficient represents the mean decrease of total surface for every additional unit in average fragmentation. If the average fragmentation increases by 1 unit, the average total surface decreases by 95.7414 (holding all the other independent variables constant).

The average parent daughter ratio coefficient in the regression equation is 54.8073. This coefficient represents the mean increase of total surface for every additional unit in average parent daughter ratio. If the average parent daughter ratio increases by 1 unit, the average total surface increases by 54.8073 (holding all the other independent variables constant).

The max branch order coefficient in the regression equation is -26.1737. This coefficient represents the mean decrease of total surface for every additional unit in max branch order. If the max branch order increases by 1 unit, the average total surface decreases by 26.1737 (holding all the other independent variables constant).

The max euclidean distance coefficient in the regression equation is 636.0267. This coefficient represents the mean increase of total surface for every additional unit in max euclidean distance. If the max euclidean distance increases by 1 unit, the average total surface increases by 636.0267 (holding all the other independent variables constant).

The max path distance coefficient in the regression equation is -346.6786. This coefficient represents the mean decrease of total surface for every additional unit in max path distance. If the max path distance increases by 1 unit, the average total surface decreases by 346.6786 (holding all the other independent variables constant).

The number bifurcations distance coefficient in the regression equation is -3895.7825. This coefficient represents the mean decrease of total surface for every additional unit in number bifurcations. If the number bifurcations increases by 1 unit, the average total surface decreases by 3895.7825 (holding all the other independent variables constant).

The number branches distance coefficient in the regression equation is 1985.8946. This coefficient represents the mean increase of total surface for every additional unit in number branches. If the number branches increases by 1 unit, the average total surface increases by 1985.8946 (holding all the other independent variables constant).

The number nodes coefficient in the regression equation is -813.9945. This coefficient represents the mean decrease of total surface for every additional unit in number nodes. If the number nodes increases by 1 unit, the average total surface decreases by 813.9945 (holding all the other independent variables constant).

The number stems coefficient in the regression equation is -458.2155. This coefficient represents the mean decrease of total surface for every additional unit in number stems. If the number stems increases by 1 unit, the average total surface decreases by 458.2155 (holding all the other independent variables constant).

The number tips distance coefficient in the regression equation is 1897.7216. This coefficient represents the mean increase of total surface for every additional unit in number tips. If the number tips increases by 1 unit, the average total surface increases by 1897.7216 (holding all the other independent variables constant).

When all the predictor measurements are equal to zero, the average total surface is 5705.7931. This does not make sense to interpret since it is not possible for neurons measurements to be zero. However, we still

need to keep the intercept term in the model in order to use it to make predictions. The intercept just doesn't have any meaningful interpretation for this model.

Second of all, I observe the p-value then comparing it with significance level of 5%. We can see that intercept, average diameter, overall depth, overall height, soma surface, total length, total volume, average parent daughter ratio, max euclidean distance, and number nodes have p-value < 0.05. Thus, these measurements are considered as significant. This also indicates that relations exist between them with the total surface.

Third of all, the AIC and BIC from this fitted model are 1.076×10^4 and 1.085×10^4 respectively. Later, we will compare this value with the forwardAIC and forwardBIC results. Lower AIC and BIC scores are better. AIC and BIC penalizes models that use more parameters. So if two models explain the same amount of variation, the one with fewer parameters will have a lower AIC and BIC score and will be the better-fit model³.

Lastly, I observe the adjusted R^2 in this fitted model which is 0.991. Thus, 99.1% of the variation in total surface can be explained by all the features in this data set.

(c) Perform a forward selection Akaike Information Criterion (AIC) regression. Examine the selected model and discuss your findings in relation to the model fitted in part (b)

From the forwardAIC result below, we can see that the important measurements that lower the AIC are: total volume, total length, average diameter, overall depth, number stems, number nodes, soma surface, average parent daughter ratio, average contraction, number bifurcations, average fragmentation, and overall width.

Based on the p-value, at significance level 5%, we see that the intercept, total volume, total length, average diameter, overall depth, number stems, number average parent daughter ratio, number bifurcations, average fragmentation are significant (p-value < 0.05). Indicating that relation exist between them with the total surface.

With this new AIC model, we have the same adjusted R^2 with the previous model that includes all variables, which is 0.991. Thus, 99.1% of the variation in total surface can be explained by total volume, total length, average diameter, overall depth, number stems, number nodes, soma surface, average parent daughter ratio, average contraction, number bifurcations, average fragmentation, and overall width.

We can see from the fitted model using forwardAIC model that the AIC value is 1.077×10^4 which (increased slightly from the previous model) while the BIC value is $1.083e \times 10^4$ (decreased slightly from the previous model).

The code of forwardAIC function below is from Week 9 Assessed Exercise Solution which return the column numbers of the X matrix for the model that gives the lowest AIC. The algorithm is as per below:

1. Begin with a model that contains no variable (other than the intercept). Run a linear regression and record the AIC. For now, this is our current model.
2. Find the most significant variable, i.e. the variable that lowers the AIC the most a. Run a linear regression with the current model plus one additional variable, and record the decrease in AIC. b. Repeat step 2a for each variable not included in the current model. c. Find the variable with the biggest decrease in AIC. d. Update the current model to include the variable that decreases the AIC the most.
3. If none of the variables lower the AIC then go to step 4. Otherwise repeat step 2 until adding variables no longer reduces the AIC.
4. Report your final chosen variables

In [31]:

```
def forwardAIC(X,y):

    mod = sm.OLS(y, X.iloc[:,0]).fit()
    best_aic = mod.aic

    bad_model = True
    chosen_vars = [0]
    remaining_vars = range(1,X.shape[1])

    while(bad_model):
        curr_aic = np.empty(len(remaining_vars))
        curr_aic_diff = np.empty(len(remaining_vars))

        for count, i in enumerate(remaining_vars):
            curr_vars = np.append(chosen_vars,i)
            curr_mod = sm.OLS(y, X.iloc[:,curr_vars]).fit()
            curr_aic[count] = curr_mod.aic
```

```

        curr_aic_diff[count] = curr_mod.aic - best_aic
    if len(remaining_vars)==0:
        bad_model=False
    elif np.min(curr_aic_diff)>0:
        bad_model = False
    else:
        best_var = remaining_vars[np.argmin(curr_aic_diff)]
        best_aic = curr_aic[np.argmin(curr_aic_diff)]
        chosen_vars = np.append(chosen_vars,best_var)
        remaining_vars = [x for x in remaining_vars if x != best_var]
    return chosen_vars

```

In [32]:

```

ans_AIC = forwardAIC(X,y)
print(ans_AIC)
print(X.columns[ans_AIC])

```

```

[ 0  7  6  1  2 18 17  5 11  9 15 10  4]
Index(['intercept', 'total_volume', 'total_length', 'average_diameter',
      'overall_depth', 'number_stems', 'number_nodes', 'soma_surface',
      'average_parent_daughter_ratio', 'average_contraction',
      'number_bifurcations', 'average_fragmentation', 'overall_width'],
      dtype='object')

```

In [33]:

```

#Fit the linear regression using forwardAIC measurements
mod = sm.OLS(y, X.iloc[:,ans_AIC])
res = mod.fit()
print(res.summary())

```

```

                    OLS Regression Results
=====
Dep. Variable:      total_surface      R-squared:      0.991
Model:              OLS               Adj. R-squared:  0.991
Method:             Least Squares     F-statistic:    6583.
Date:               Sun, 11 Dec 2022   Prob (F-statistic): 0.00
Time:               16:26:06           Log-Likelihood: -5371.2
No. Observations:   694               AIC:            1.077e+04
Df Residuals:       681               BIC:            1.083e+04
Df Model:           12
Covariance Type:    nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
intercept	5705.7931	21.302	267.851	0.000	5663.967	5747.619
total_volume	3006.6073	48.140	62.455	0.000	2912.087	3101.128
total_length	3738.2780	406.017	9.207	0.000	2941.083	4535.473
average_diameter	691.1032	35.595	19.416	0.000	621.215	760.992
overall_depth	-135.2543	29.395	-4.601	0.000	-192.970	-77.538
number_stems	-56.7562	24.052	-2.360	0.019	-103.980	-9.532
number_nodes	-598.5320	405.474	-1.476	0.140	-1394.662	197.598
soma_surface	48.4343	28.267	1.713	0.087	-7.067	103.935
average_parent_daughter_ratio	47.3583	23.019	2.057	0.040	2.161	92.556
average_contraction	38.2642	23.536	1.626	0.104	-7.948	84.476
number_bifurcations	-212.3105	64.774	-3.278	0.001	-339.492	-85.129
average_fragmentation	-148.1038	44.649	-3.317	0.001	-235.770	-60.437
overall_width	64.7112	35.276	1.834	0.067	-4.551	133.973

```

=====
Omnibus:      603.992   Durbin-Watson:      1.892
Prob(Omnibus): 0.000   Jarque-Bera (JB):    32666.481
Skew:         -3.553   Prob(JB):            0.00
Kurtosis:     35.851   Cond. No.            60.8
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

(d) Perform a forward selection Bayes Information Criterion (BIC) regression. Examine the selected model and discuss your findings in relation to the models fitted in part (b) and (c).

From the forwardBIC result below, we can see that the important measurements that lower the BIC are: total volume, total length, average_diameter, overall_depth, number stems, number nodes.

Based on the p-value, at significance level 5%, we see that the intercept, total volume, total length, average diameter, overall depth, and number stems are significant (p-value < 0.05). Indicating that relation exist between these variables with the total surface.

With this new BIC model, we have the same adjusted R^2 with the model that includes all variables and the AIC model, which is 0.991. Thus, 99.1% of the variation in total surface can be explained by all the variables in this data set.

We can see from the fitted model using forwardBIC model that the AIC value is 1.078×10^4 which (increased slightly from the previous model) while the BIC value is $1.081e \times 10^4$ (decreased slightly from the previous model).

Here I use forwardBIC function that has the same algorithm like forwardAIC from Week 9 Assessed Exercise Solution. This return the column numbers of the X matrix for the model that gives the lowest BIC. The algorithm is as per below:

1. Begin with a model that contains no variable (other than the intercept). Run a linear regression and record the BIC. For now, this is our current model.
2. Find the most significant variable, i.e. the variable that lowers the BIC the most a. Run a linear regression with the current model plus one additional variable, and record the decrease in BIC. b. Repeat step 2a for each variable not included in the current model. c. Find the variable with the biggest decrease in BIC. d. Update the current model to include the variable that decreases the BIC the most.
3. If none of the variables lower the BIC then go to step 4. Otherwise repeat step 2 until adding variables no longer reduces the AIC.
4. Report your final chosen variables

```
In [34]: def forwardBIC(X,y):

    mod = sm.OLS(y, X.iloc[:,0]).fit()
    best_bic = mod.bic

    bad_model = True
    chosen_vars = [0]
    remaining_vars = range(1,X.shape[1])
    while(bad_model):
        curr_bic = np.empty(len(remaining_vars))
        curr_bic_diff = np.empty(len(remaining_vars))
        for count, i in enumerate(remaining_vars):
            curr_vars = np.append(chosen_vars,i)
            curr_mod = sm.OLS(y, X.iloc[:,curr_vars]).fit()
            curr_bic[count] = curr_mod.bic
            curr_bic_diff[count] = curr_mod.bic - best_bic
        if len(remaining_vars)==0:
            bad_model=False
        elif np.min(curr_bic_diff)>0:
            bad_model = False
        else:
            best_var = remaining_vars[np.argmin(curr_bic_diff)]
            best_bic = curr_bic[np.argmin(curr_bic_diff)]
            chosen_vars = np.append(chosen_vars,best_var)
            remaining_vars = [x for x in remaining_vars if x != best_var]
    return chosen_vars
```

```
In [35]: ans_BIC = forwardBIC(X,y)
print(ans_BIC)
print(X.columns[ans_BIC])

[ 0  7  6  1  2 18 17]
Index(['intercept', 'total_volume', 'total_length', 'average_diameter',
      'overall_depth', 'number_stems', 'number_nodes'],
      dtype='object')
```

```
In [36]: #Fit the linear regression using forwardBIC measurements
mod = sm.OLS(y, X.iloc[:,ans_BIC])
res = mod.fit()
print(res.summary())
```

OLS Regression Results			
Dep. Variable:	total_surface	R-squared:	0.991
Model:	OLS	Adj. R-squared:	0.991
Method:	Least Squares	F-statistic:	1.279e+04
Date:	Sun, 11 Dec 2022	Prob (F-statistic):	0.00
Time:	16:26:06	Log-Likelihood:	-5384.2
No. Observations:	694	AIC:	1.078e+04
Df Residuals:	687	BIC:	1.081e+04

Df Model:	6					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
intercept	5705.7931	21.609	264.047	0.000	5663.366	5748.221
total_volume	3030.2338	47.786	63.412	0.000	2936.409	3124.059
total_length	4017.9089	397.912	10.097	0.000	3236.639	4799.179
average_diameter	709.4829	31.482	22.536	0.000	647.670	771.296
overall_depth	-168.0593	27.537	-6.103	0.000	-222.126	-113.992
number_stems	-67.3998	23.011	-2.929	0.004	-112.580	-22.220
number_nodes	-1057.0739	391.404	-2.701	0.007	-1825.565	-288.583
Omnibus:	577.310	Durbin-Watson:		1.880		
Prob(Omnibus):	0.000	Jarque-Bera (JB):		28055.491		
Skew:	-3.339	Prob(JB):		0.00		
Kurtosis:	33.424	Cond. No.		47.5		

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

(e) Explain how using BIC for model selection differs from using AIC

AIC is an estimate of a constant plus the relative distance between the unknown true likelihood function of the data and the fitted likelihood function of the model, whereas BIC is an estimate of a function of the posterior probability of a model being true, under a certain Bayesian setup.

They consist of a goodness-of-fit term plus a penalty to control over-fitting, and provide a standardized way to balance between having enough parameters to adequately model the relationships among variables in the population (sensitivity) and not over-fitting a model or suggesting nonexistent relationships (specificity). Hence the smaller the AIC/BIC, the better the model.

AIC and BIC difference is the size of the penalty. BIC tends to be more restrictive when it comes to penalizing model complexity even though it is sometimes preferred over AIC for some of the reasons below:

- 1) BIC is consistent. A consistent selector is one which will select the true model with probability approaching 100% as n tends to infinity. AIC is not consistent because it has a non-vanishing chance of choosing an unnecessarily complex model as n becomes large.
- 2) BIC considers Type I and Type II errors to be about equally undesirable, while AIC considers Type II errors to be more undesirable than Type I errors unless n is very small. In classical hypothesis testing, over-fitting (Type I errors) are considered worse than Type II errors, whereas for prediction, a Type II error can be harmful as well.²

Question 5 - Random Forest Regression to Predict the Total Surface Area of a Neuron

(a) Split the data into appropriate training and test sets.

We have 694 observations so will split the data into training and test sets of size 521 and 173 respectively. Importantly, the data for each set must be selected randomly and each observation must only appear in one of the sets.

In [37]:

```
#Setting the random seed
np.random.seed(123)

#Here I divide the train size with 75% for training purposes
#and the remaining 25% for the test data
train_size = 521

#Create new data for X to be used for random forest regression
#because we want to use the non-standardised dataset, X_rfr.
X_rfr = df_neurons_all.drop(['id', 'total_surface'], axis=1)
train_select = np.random.permutation(range(len(y)))
X_train = X_rfr.iloc[train_select[:train_size],:].reset_index(drop=True)
X_test = X_rfr.iloc[train_select[train_size:],:].reset_index(drop=True)

#Use the y data that we have already created in above part
#because it is not standardized
y_train = y[train_select[:train_size]].reset_index(drop=True)
y_test = y[train_select[train_size:]].reset_index(drop=True)

#reset_index drops the original row index and resets it,
#such that the first row has index 0, second has index 1, and so on
```


(b) Fit a random forest regression model with 10 trees using the training data. Include the argument `random state=101` in the random forest regression function to ensure reproducible results. Determine which variables are most important in predicting the total surface area of a neuron. Discuss your findings in relation to the linear models fit in question 4.

Here, I fit a random forest regression model to the training data with 10 trees and random state = 101 in the random forest regression function to ensure reproducible results. After that, I use feature importances attribute from the Random Forest Regressor to obtain the most important variable in predicting the total surface area of a neuron.

We can see from the result below that by using random forest regression, the four variable that are important are total volume, number bifurcations, total length, and number nodes. We get the same result for variable Total Volume in AIC, BIC, and Random Forest Regression model that this variable came out as the most important variable in predicting total surface.

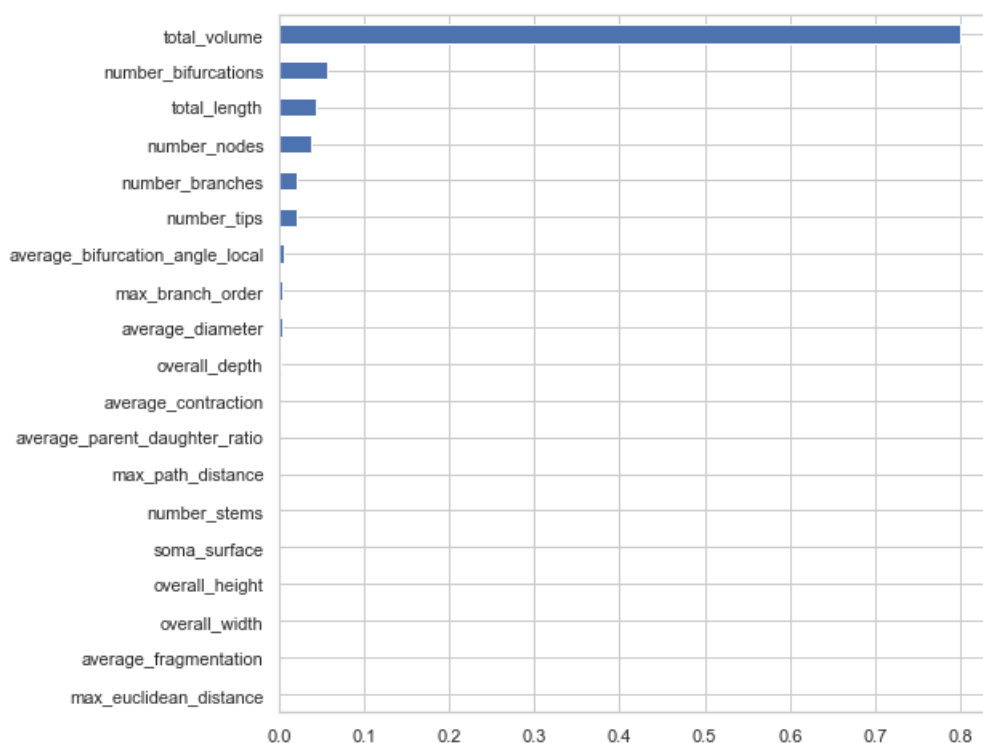
```
In [38]: #First, I import RandomForestRegressor
#function from sklearn.ensemble
from sklearn.ensemble import RandomForestRegressor

#Initializing the Random Forest Regression model
#with 10 decision trees and using random state 101
#to ensure reproducible results
regressor = RandomForestRegressor(n_estimators = 10, random_state = 101)

#Fitting Random Forest Regression to the training dataset
reg = regressor.fit(X_train,y_train)
```

By using random forest regression model, the variables that are most important in predicting the total surface area of a neuron is total volume.

```
In [39]: ##Use feature_importances_ attribute from the Random Forest Regressor
##to obtain the most important variable in
##predicting the total surface area of a neuron
plt.figure(figsize=(8,8))
important_feature = pd.Series(reg.feature_importances_,index=X_train.columns).sort_values(ascending=False)
ax = important_feature.plot(kind='barh')
plt.show()
```

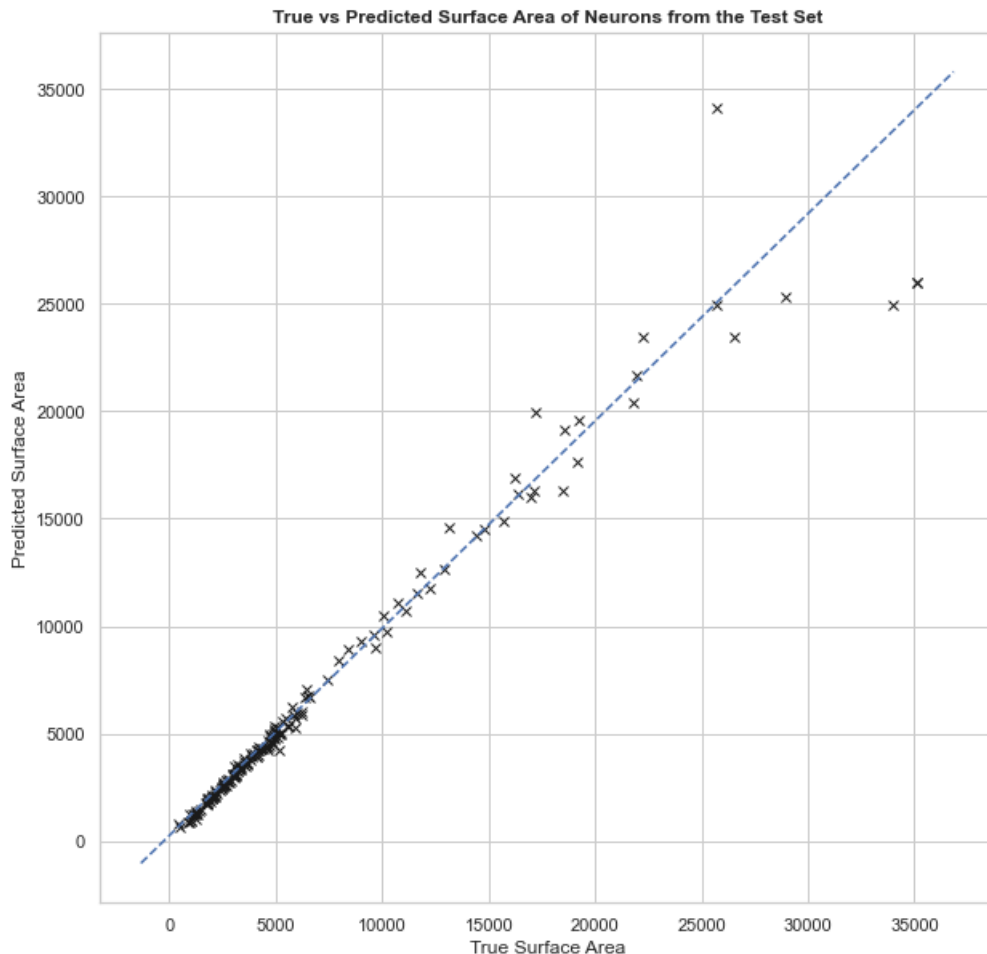


(c) Use the random forest regression model to predict the total surface area of a neuron for the test set. Create a scatter plot of the true surface area of a neuron versus the predicted surface area. Interpret your plot.

From the scatterplot below, we can see that the true surface area of a neuron and the predicted surface area have a strong positive linear relationship between each other indicating that when true surface area increase, the predicted surface area increases in a similar fashion.

```
In [40]: #Use predict function from the Random Forest Regressor
#to predict the total surface area of a neuron from the test set
prediction = reg.predict(X_test)

#Then, I create a scatterplot for the true surface area
#of a neuron (x-axis) versus the predicted surface area (y-axis)
plt.figure(figsize=(10,10))
plt.plot(y_test,prediction,'kx')
plt.plot(plt.xlim(), plt.ylim(), ls="--")
plt.xlabel('True Surface Area')
plt.ylabel('Predicted Surface Area')
plt.title('True vs Predicted Surface Area of Neurons from the Test Set',fontweight="bold")
plt.show()
```



(d) Assess the performance of a random forest regression model with 5, 10, 20, 50, 100, 200, 500 and 1000 trees in predicting the total surface area of a neuron. You should repeat the model fit and prediction 30 times for each number of trees, using a different random state for each repeat. Create a plot of the model performance as a function of the number of trees (use a log axis for the number of trees). The plot should show the mean and standard error of the performance metric for each number of trees. Discuss your findings.

We may see in the plot and result below that the MSE decreases as the number of tree increases. We obtain the lowest MSE of mean and standard deviation when the number of trees are 1000.

Generally, more trees is equivalent to more features/parameters in our model. A higher number of features always reduces training error. This is simply due to the fact that if those additional features are unhelpful, then those features will not be used and the training error will at least remain the same as the model with fewer features⁴.

```
In [41]: from sklearn.metrics import mean_squared_error

#Create list that consist of 5, 10, 20, 50, 100, 200, 500 and 1000 trees
trees = [5, 10, 20, 50, 100, 200, 500, 1000]

#Create random state range value
random_state = range(1,31)

#Initializing MSE list and dictionary
mse_list = []
mse_mean_dict = {}
```

```

mse_se_dict = {}

#Repeat the model fit and prediction 30 times for each number of trees,
#using a different random state for each repeat
for tree in trees:
    for value in random_state:
        rfr = RandomForestRegressor(n_estimators = tree, random_state = value)
        rfr.fit(X_train, y_train)
        rfr_test_pred = rfr.predict(X_test)
        mse = mean_squared_error(y_test, rfr_test_pred)
        mse_list.append(mse)
        mse_array = np.array(mse_list)

        #Calculate MSE Mean
        mse_mean = np.mean(mse_array)

        #Calculate MSE Standard Error. The formula of standard error
        #standard deviation divided by the sample size's square root
        mse_std = np.std(mse_array)
        mse_se = mse_std/np.sqrt(len(random_state))
        mse_mean_dict[str(tree)] = mse_mean
        mse_se_dict[str(tree)] = mse_se

```

In [42]:

```

#Creating new dataset consisting of tree, mse mean, and mse standard error using below code:
mse_mean_df = pd.DataFrame(list(mse_mean_dict.items()))
mse_mean_df.rename(columns = {0:'tree', 1:"mse_mean"}, inplace = True)

mse_se_df = pd.DataFrame(list(mse_se_dict.items()))
mse_se_df.rename(columns = {0:'tree', 1:"mse_se"}, inplace = True)

mse_df = pd.merge(mse_mean_df,mse_se_df,on=['tree'],how = 'inner')
print(mse_df)
print(mse_df.info())

```

```

   tree  mse_mean  mse_se
0     5  1.911145e+06  133440.742507
1    10  1.788562e+06  131286.327803
2    20  1.667164e+06  121332.733373
3    50  1.566631e+06  112777.309079
4   100  1.507238e+06  104609.862007
5   200  1.469564e+06   97376.502435
6   500  1.441122e+06   91216.568383
7  1000  1.420199e+06   86013.560322
<class 'pandas.core.frame.DataFrame'>
Int64Index: 8 entries, 0 to 7
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   tree        8 non-null         object
1   mse_mean    8 non-null         float64
2   mse_se      8 non-null         float64
dtypes: float64(2), object(1)
memory usage: 256.0+ bytes
None

```

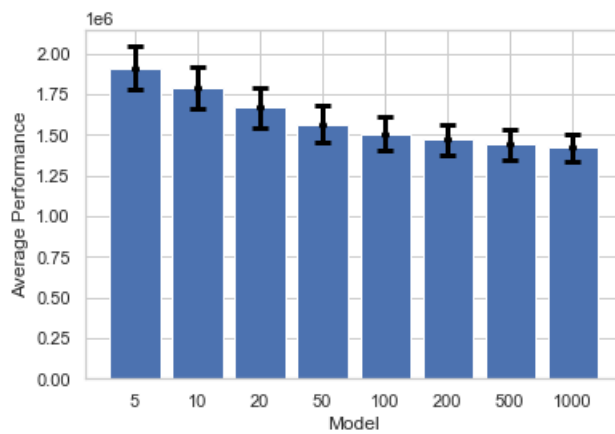
In [43]:

```

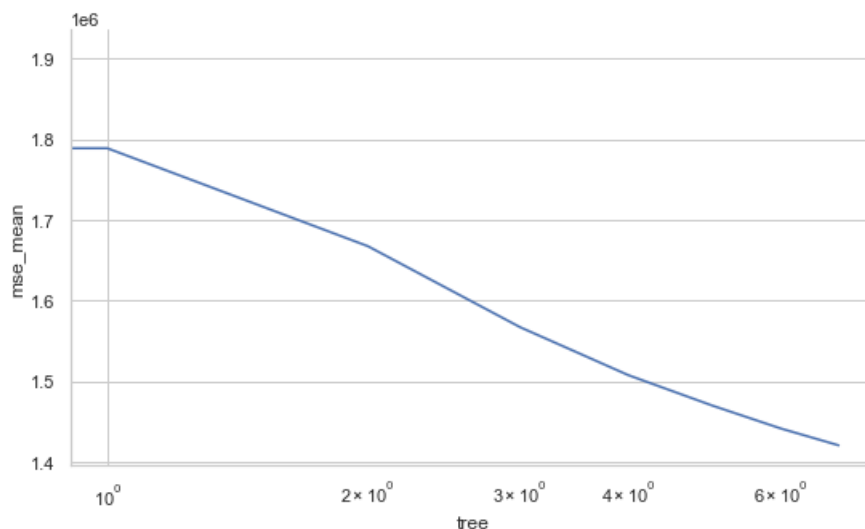
#Create the error bar
plt.errorbar(mse_df['tree'],
             mse_df['mse_mean'],
             yerr=mse_df['mse_se'],
             fmt='o', color='Black',
             elinewidth=3,capthick=3,
             errorevery=1,
             alpha=1, ms=4,
             capsize = 5)

#Create the bar plot
plt.bar(mse_df['tree'],
        mse_df['mse_mean'],
        tick_label = mse_df['tree'])
plt.xlabel('Model')
plt.ylabel('Average Performance')
plt.show()

```



```
In [44]: #Create lineplot with error band based on the standard error
fig, ax = plt.subplots(figsize=(9,5))
sns.lineplot(data=mse_df,
              x="tree",
              y="mse_mean",
              ci='mse_se')
plt.xscale('log')
sns.despine()
plt.show()
```



(e) Explain the rationale for fitting the model multiple time with different random states.

Fitting the model multiple times with different random states can increase the model performance. Sometimes, we can get significant performance improvement for our model by running it multiple times with different random states. This is because random state is also a hyperparameter and we can tune it to get better results⁵.

Question 6 - Clustering algorithms to identify different neuron types

(a) Perform a k-means cluster analysis, using the morphological measurements as the features. Run the clustering algorithm for different numbers of clusters (integers from 1 to 10). Plot the model performance as a function of the number of clusters and identify the optimal number of clusters for this data.

K-Means is an unsupervised machine learning algorithm that groups data into k number of clusters. The number of clusters is user-defined and the algorithm will try to group the data even if this number is not optimal for the specific case. Therefore we have to come up with a technique that somehow will help us decide how many clusters we should use for the K-Means model.

The Elbow method is a very popular technique and the idea is to run k-means clustering for a range of clusters. We are calculating the sum of squared distances from each point to its assigned center (distortions). When the distortions are plotted and the plot looks like an arm then the "elbow" (the point of inflection on the curve) is the best value of k.⁶

We can see from the Elbow Method showing the optimal k below that the optimal k is 3.

```
In [45]: #First, I create a new dataframe without column "id"
df_cluster = df_neurons_all.drop(['id'],axis=1)
```

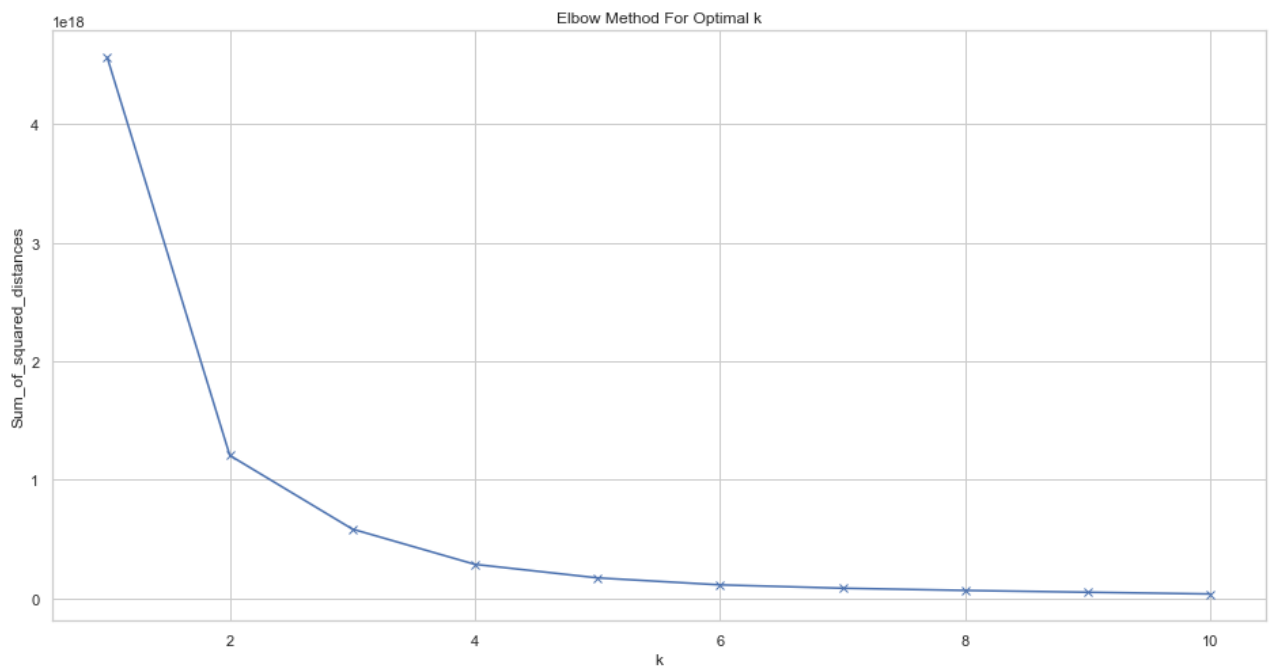
In [46]:

```
#Then, I plot the model performance as a function
#of the number of clusters and
#identify the optimal number of clusters for this data

from sklearn.cluster import KMeans

Sum_of_squared_distances = []
K = range(1,11)
for k in K:
    km = KMeans(n_clusters=k)
    km = km.fit(df_neurons_all)
    Sum_of_squared_distances.append(km.inertia_)

plt.figure(figsize=(16,8))
plt.plot(K, Sum_of_squared_distances, 'bx-')
plt.xlabel('k')
plt.ylabel('Sum_of_squared_distances')
plt.title('Elbow Method For Optimal k')
plt.show()
```



(b) Perform a k-means cluster analysis, using the optimal number of clusters (identified in part (a)), and identify the most discriminatory variables. (Hint: Create histograms for each variable, with the data separated by cluster.)

Here, I use the optimal number of clusters that I identified in part a which is 3 and identify most discriminatory variables. From the density plot that I created below, we can see that the discriminatory variables that show a clear difference between clusters (not overlapping between each other) are: total length, total surface, total volume, and number nodes.

In [47]:

```
#Fit K Means Clustering algorithm with k = 3
kmeans_model = KMeans(n_clusters = 3)
kmeans_model.fit(df_cluster)

#Creating a new column that shows which cluster
#the neuron belongs to based on K Means clustering
df_cluster['kmeans_cluster'] = kmeans_model.labels_
```

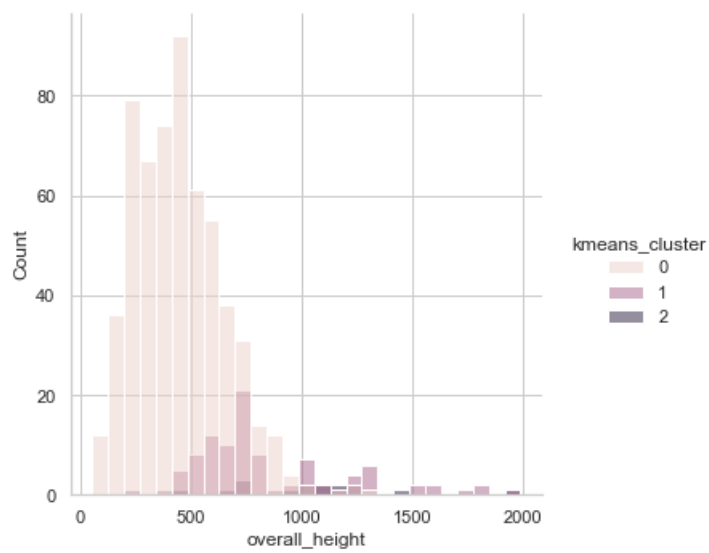
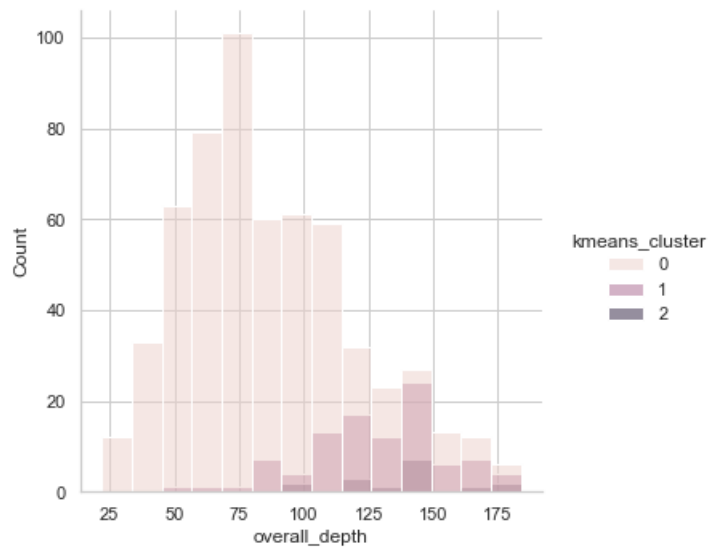
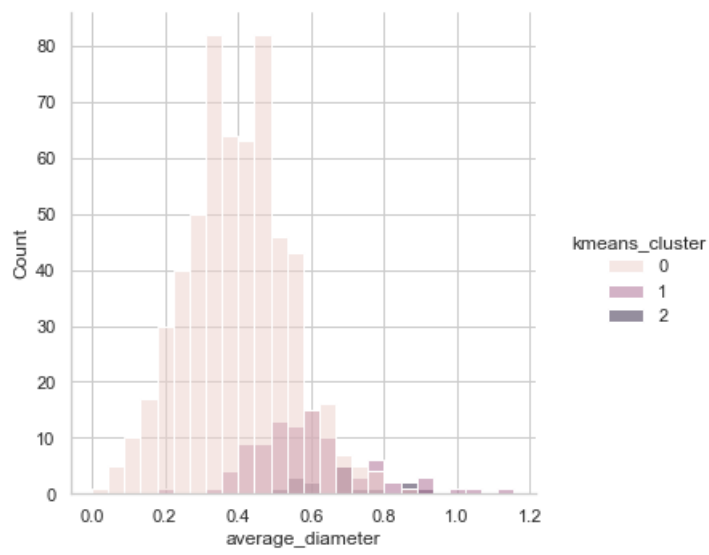
In [48]:

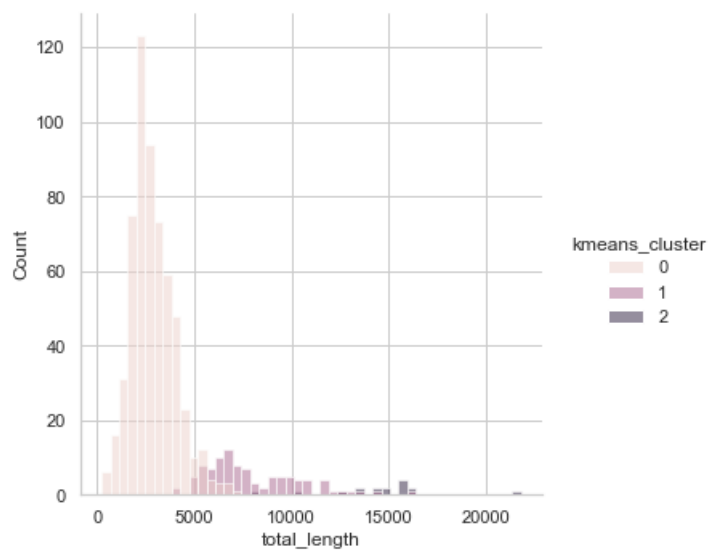
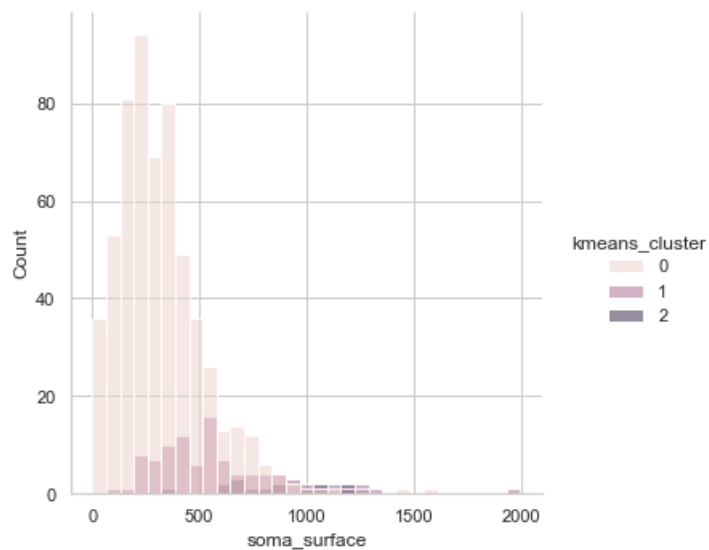
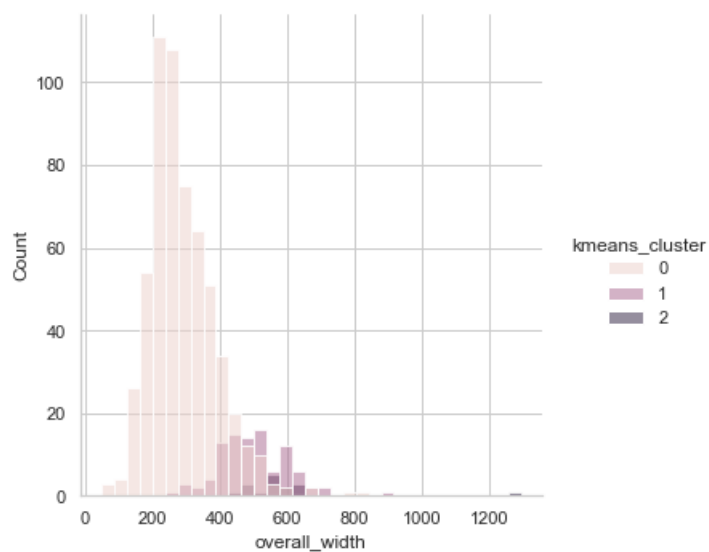
```
#First, I create a list containing all the feature names
col_list = df_cluster.columns.to_list()

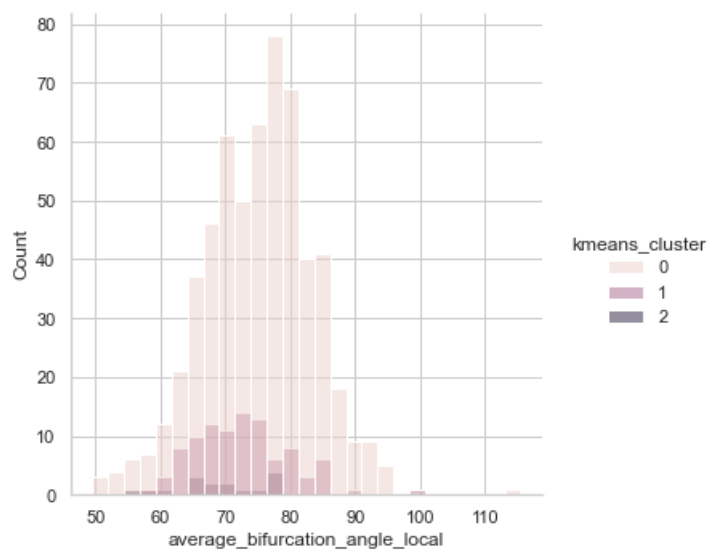
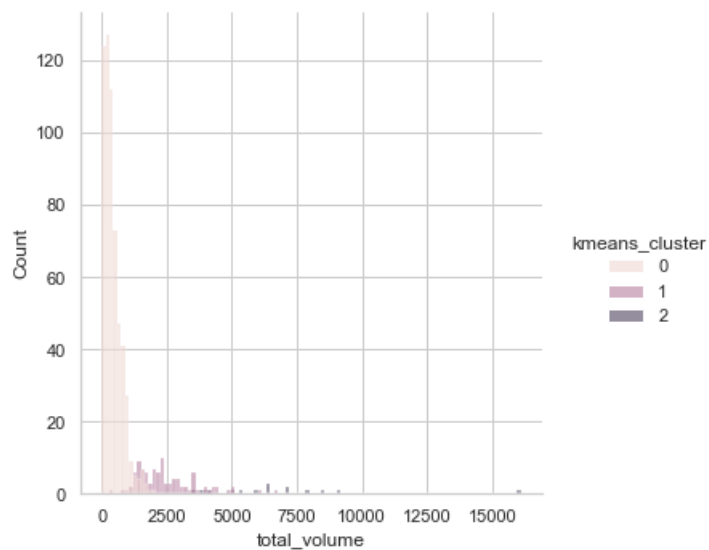
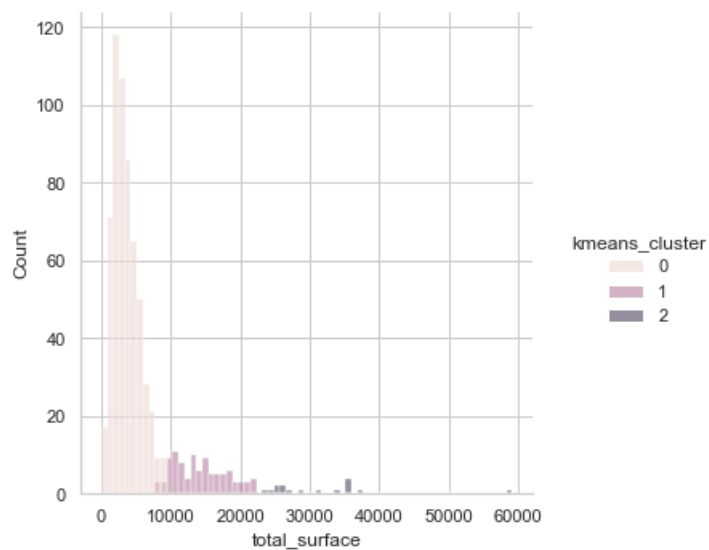
#Then, I Create histogram plot for each feature
plt.figure(figsize=(18,12))

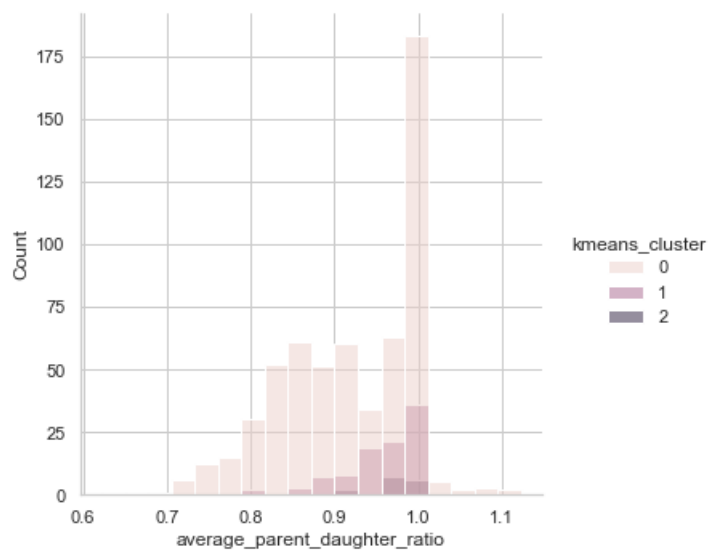
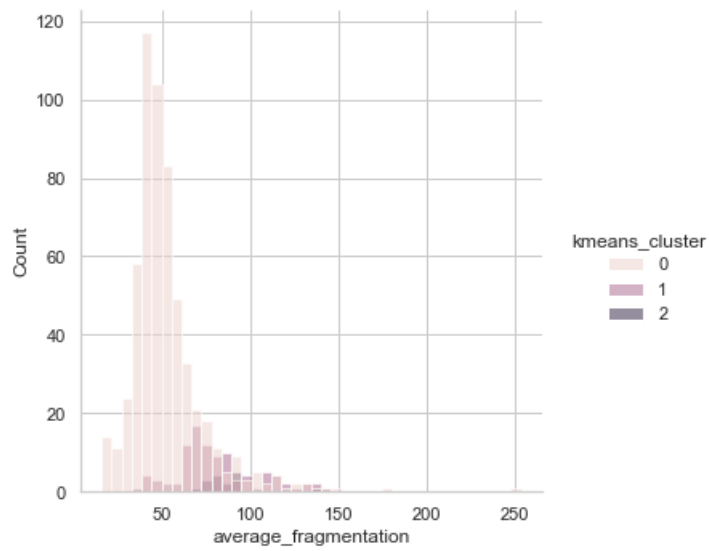
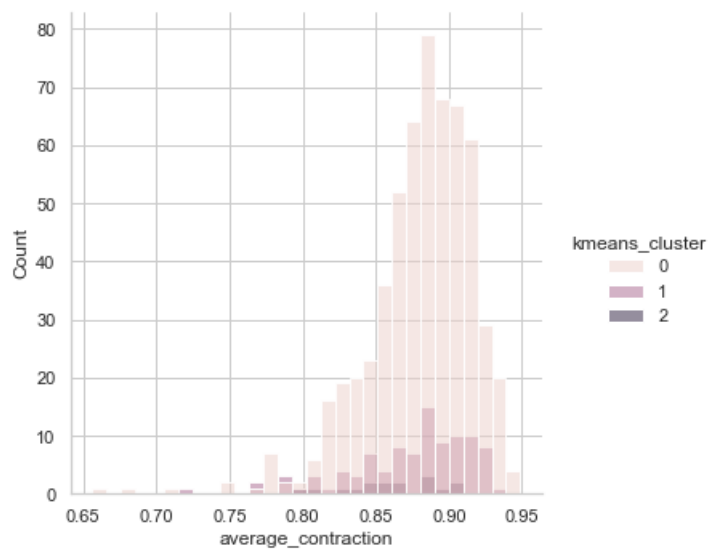
for feature in col_list:
    sns.displot(data = df_cluster, x= feature, hue="kmeans_cluster")
plt.show()
```

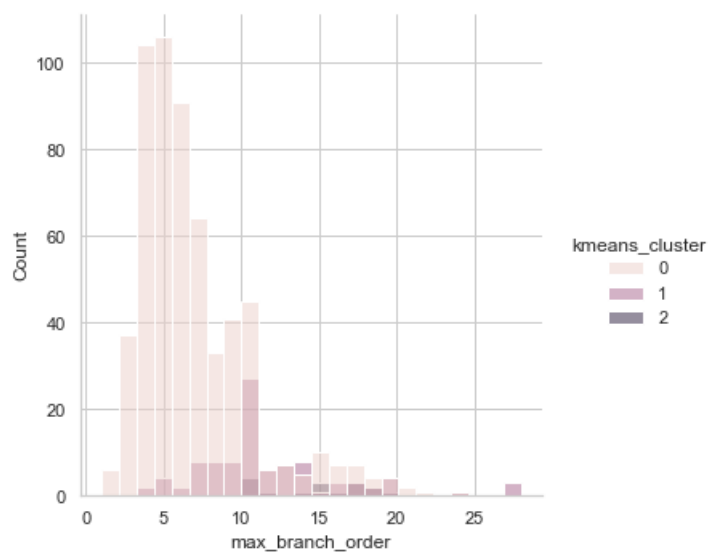
<Figure size 1296x864 with 0 Axes>

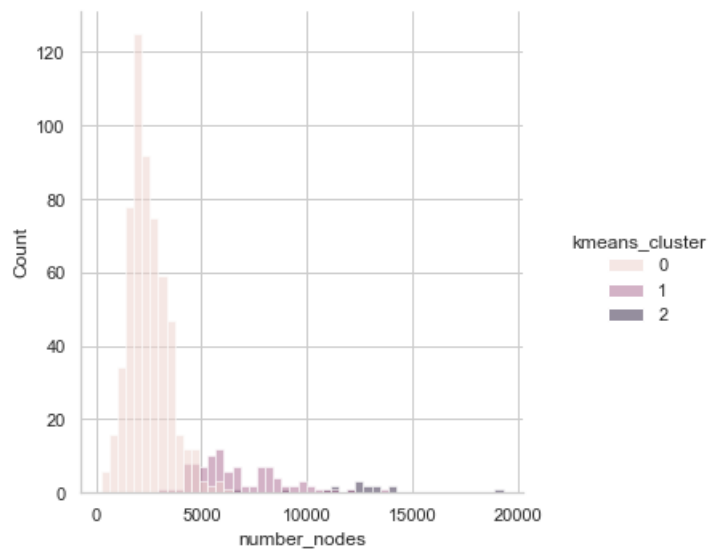
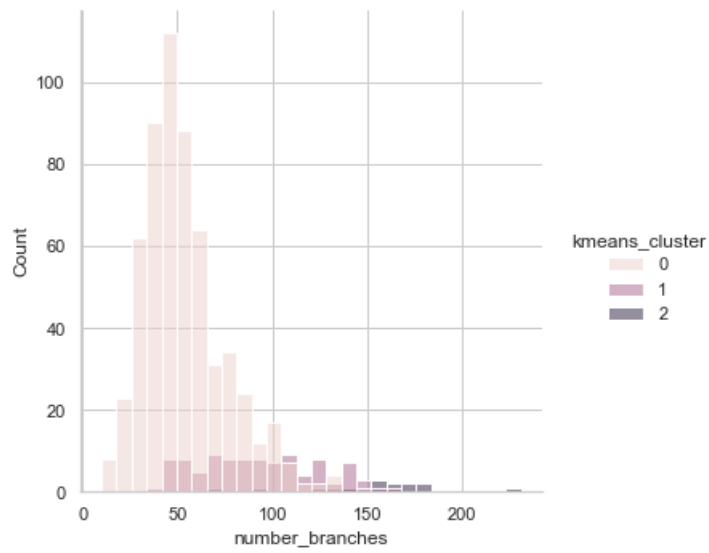
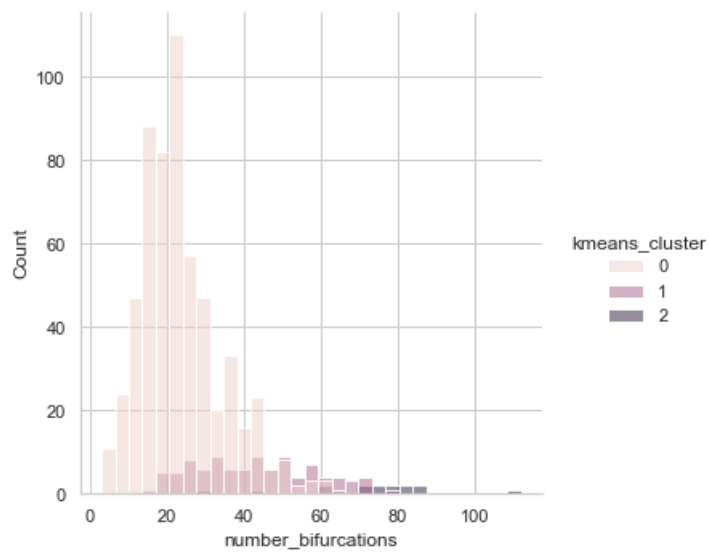


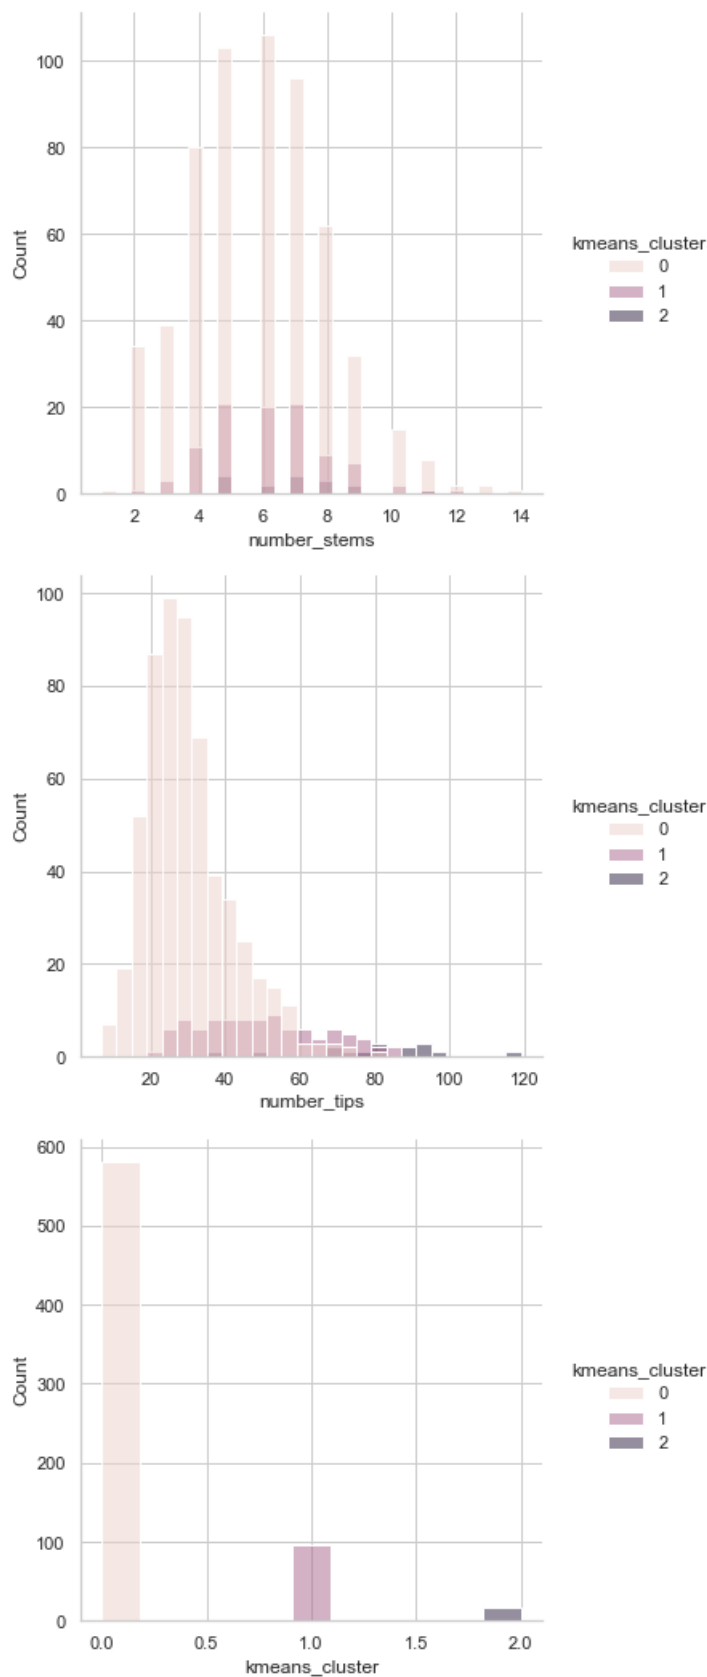












(c) Create a series of scatter plots for the most discriminatory variables, colouring the points by cluster number. Discuss your findings. Do your findings support the claim that multiple categories of neurons, with distinctly different morphological properties, are included in this dataset?

From the scatterplot below, we can see clearly that the discriminatory variables (total length, total surface, total volume, and number nodes) are mostly difference between clusters (not overlapping between each other).

```
In [49]: sns.set(style="whitegrid")

fig,axs = plt.subplots(3,2, figsize = (18,18))

#Scatterplot between number nodes and total length measurements
sns.scatterplot(data = df_cluster,
```

```

x = "total_length",
y = "total_surface",
hue = 'kmeans_cluster',
ax=axes[0, 0])

#Scatterplot between number tips and number branches measurements
sns.scatterplot(data= df_cluster,
                x="total_length",
                y="total_volume",
                hue = 'kmeans_cluster',
                ax=axes[0, 1])

#Scatterplot between number branches and number bifurcations measurements
sns.scatterplot(data=df_cluster,
                x="total_length",
                y="number_nodes",
                hue = 'kmeans_cluster',
                ax=axes[1, 0])

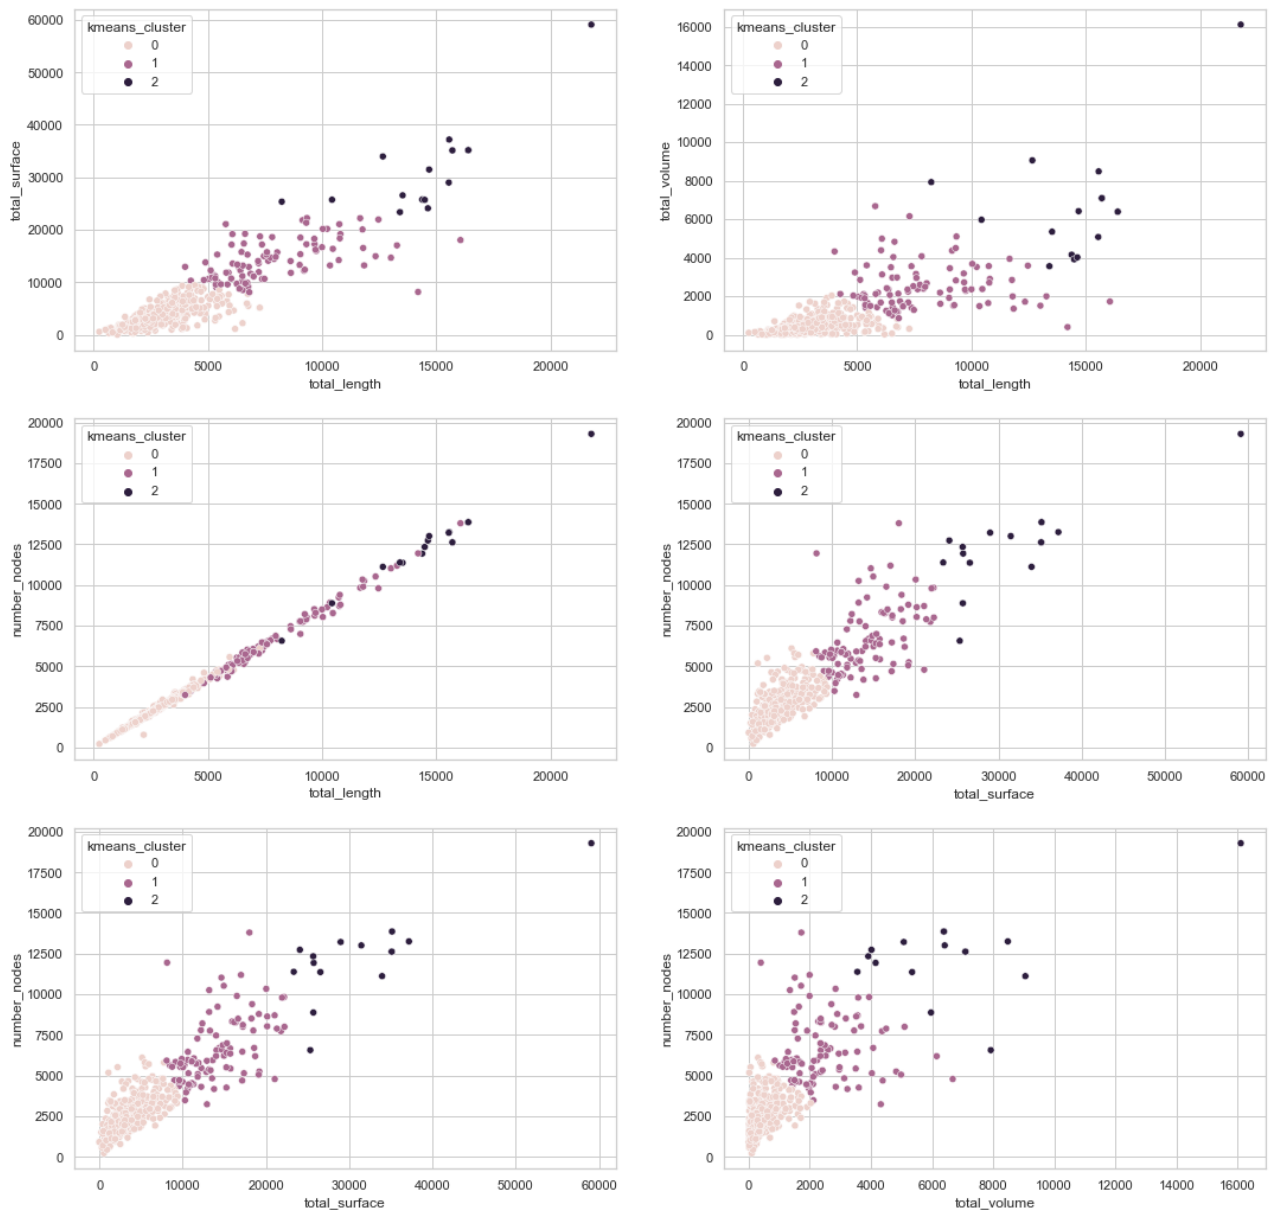
#Scatterplot between number tips and number bifurcations measurements
sns.scatterplot(data=df_cluster,
                x="total_surface",
                y="number_nodes",
                hue = 'kmeans_cluster',
                ax=axes[1, 1])

#Scatterplot between number branches and number bifurcations measurements
sns.scatterplot(data=df_cluster,
                x="total_surface",
                y="number_nodes",
                hue = 'kmeans_cluster',
                ax=axes[2, 0])

#Scatterplot between number tips and number bifurcations measurements
sns.scatterplot(data=df_cluster,
                x="total_volume",
                y="number_nodes",
                hue = 'kmeans_cluster',
                ax=axes[2, 1])

plt.show()

```



(d) Identify another clustering algorithm that may be suitable for this data. Give an overview of your chosen algorithm and discuss the type of problems it works best for. Repeat part (a)–(c) using your chosen algorithm. Discuss your results in relation to those from the k-means cluster analysis.

In this part, I use Balanced Iterative Reducing and Clustering using Hierarchies (BIRCH) method. BIRCH deals with large datasets by first generating a more compact summary that retains as much distribution information as possible, and then clustering the data summary instead of the original dataset. BIRCH actually complements other clustering algorithms by virtue of the fact that different clustering algorithms can be applied to the summary produced by BIRCH. BIRCH can only deal with metric attributes (similar to the kind of features KMEANS can handle). A metric attribute is one whose values can be represented by explicit coordinates in an Euclidean space (no categorical variables).⁷

From the histogram plot that I created below, we can see that the discriminatory variables that show a clear difference between clusters (not overlapping between each other) are: total length, total surface, total volume, and number nodes.

We obtain the same variables as the variables in the K-Means Clustering method.

```
In [50]: from sklearn.cluster import Birch

#Fit the Birch Clustering algorithm with k = 3
birch_model = Birch(n_clusters = 3)
birch_model.fit(df_cluster)

#Creating a new column that shows which cluster
#the neuron belongs to based on K Means clustering
df_cluster['birch_cluster'] = birch_model.labels_
```

```
In [51]: #Create histogram plot for each feature
```

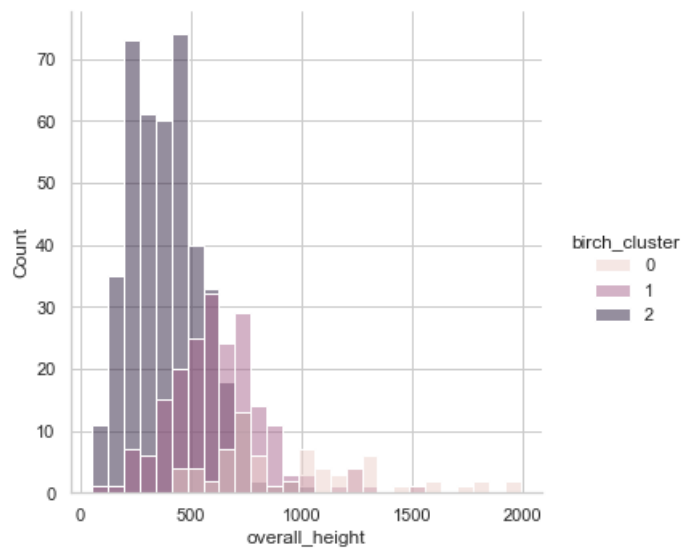
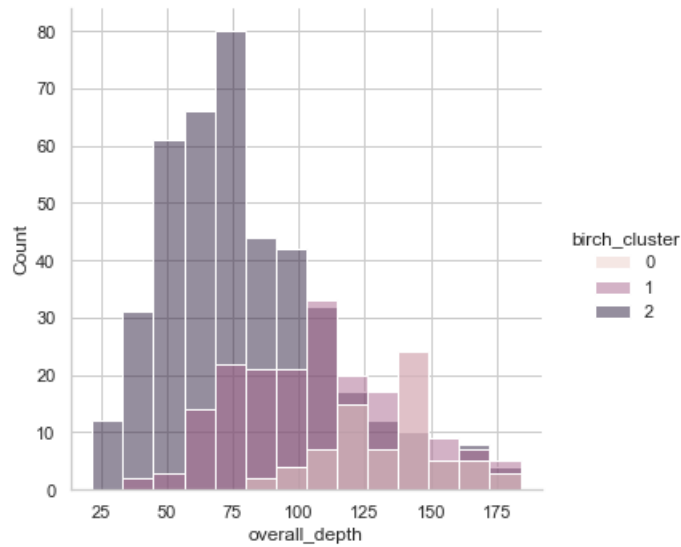
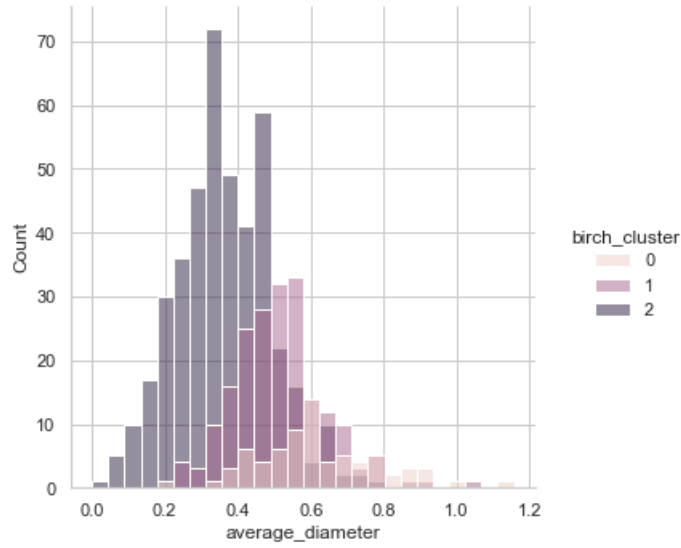


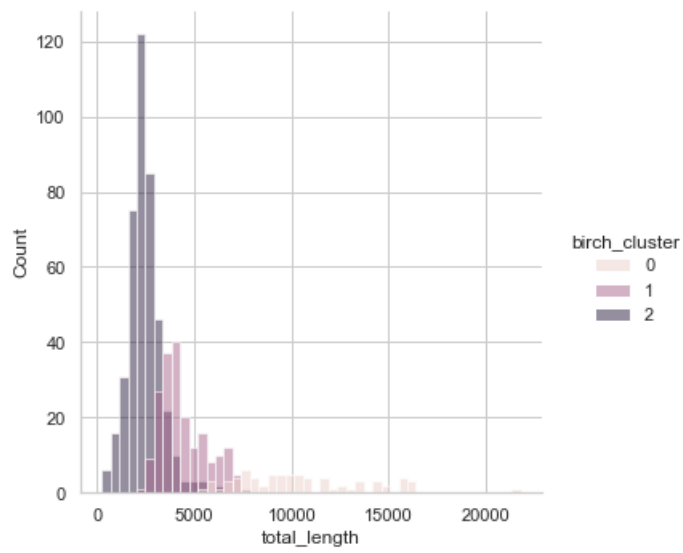
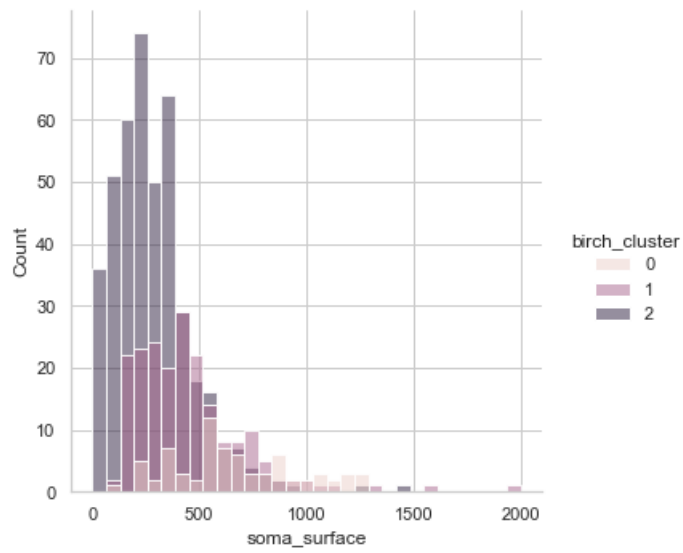
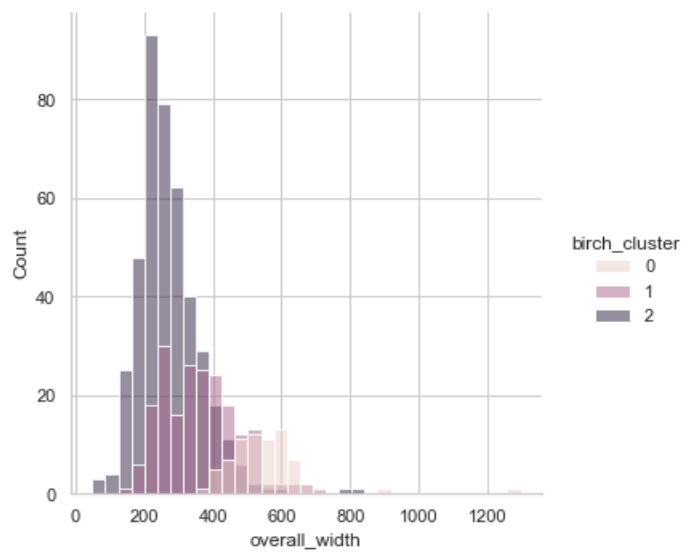
```
plt.figure(figsize=(18,12))

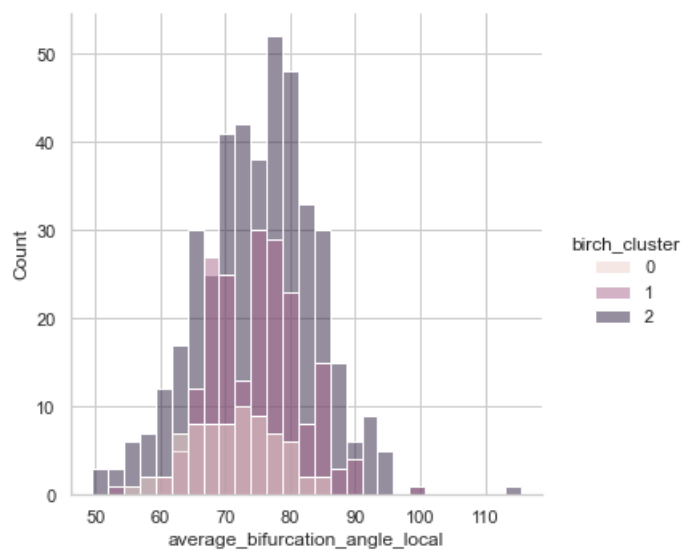
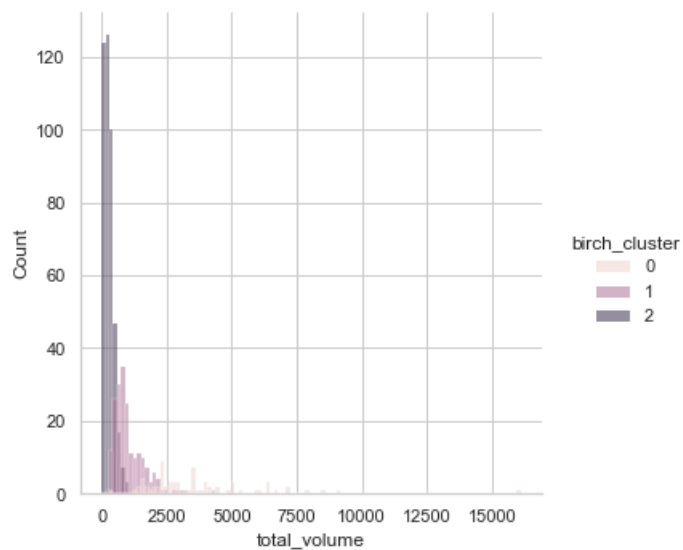
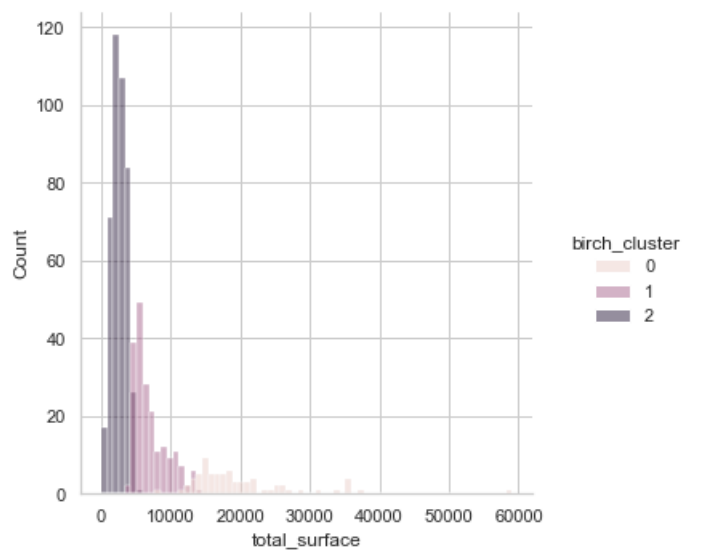
for feature in col_list:
    sns.displot(data = df_cluster,
                x= feature,
                hue="birch_cluster")

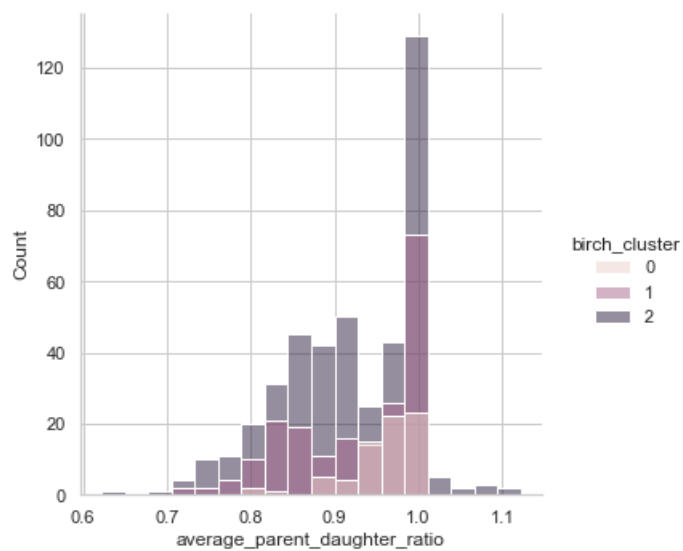
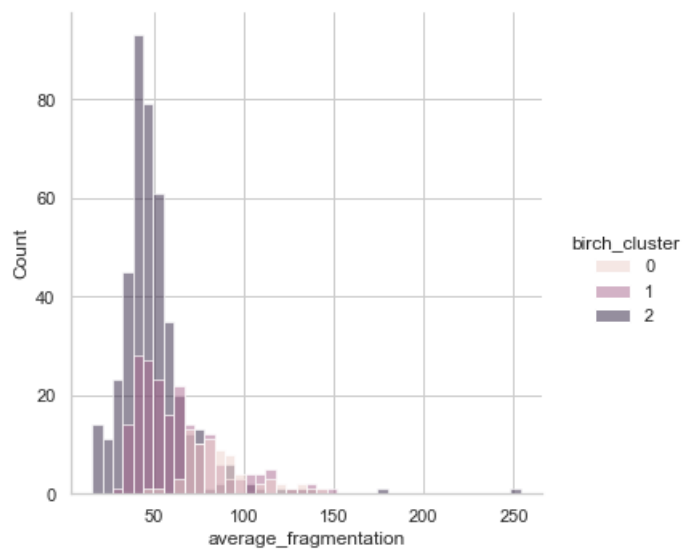
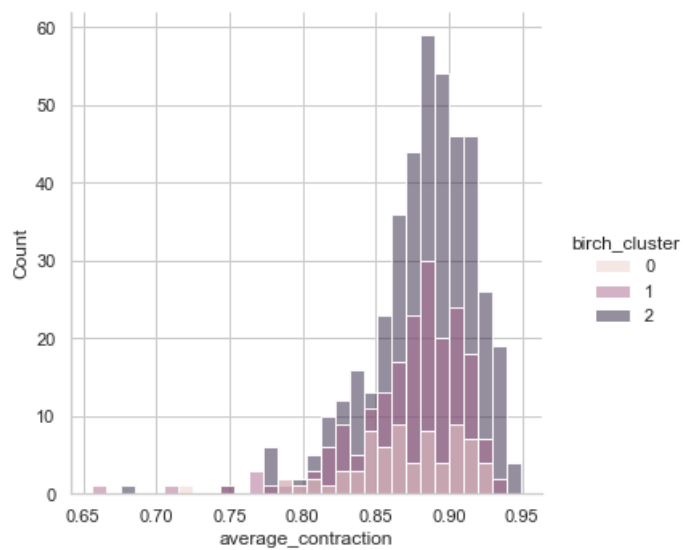
plt.show()
```

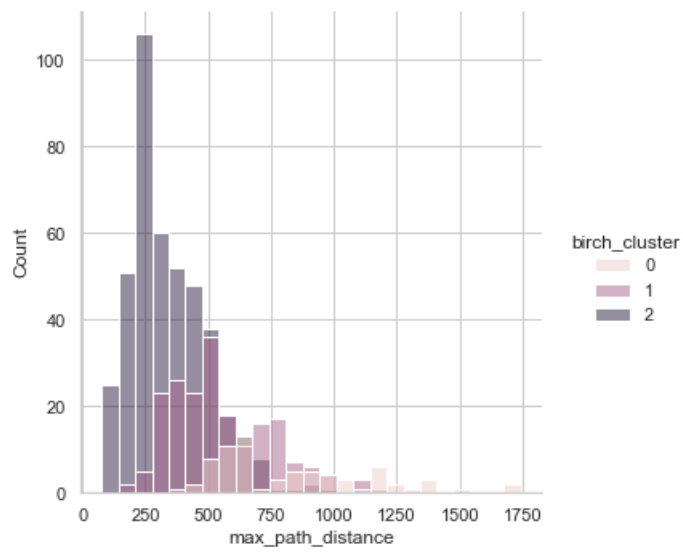
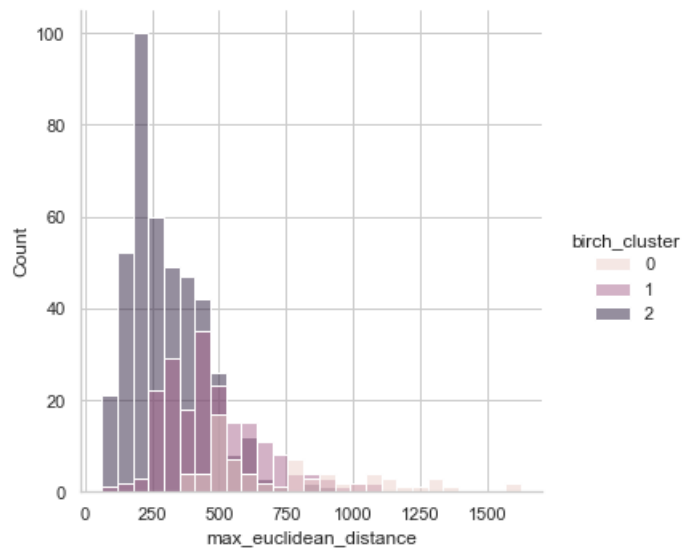
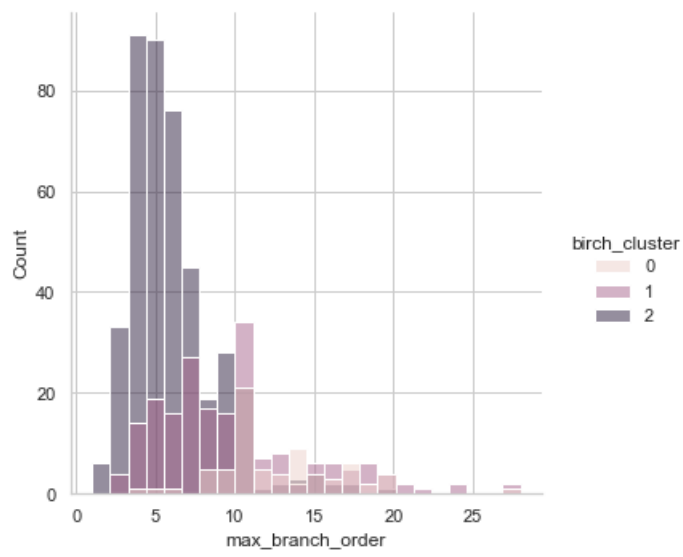
<Figure size 1296x864 with 0 Axes>

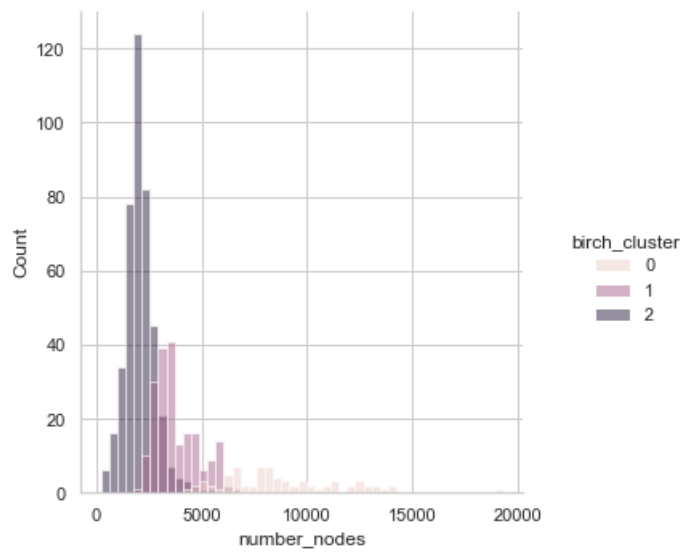
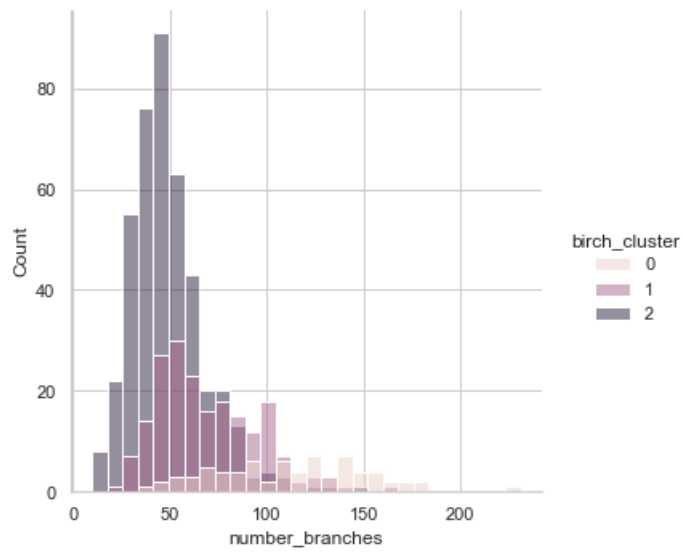
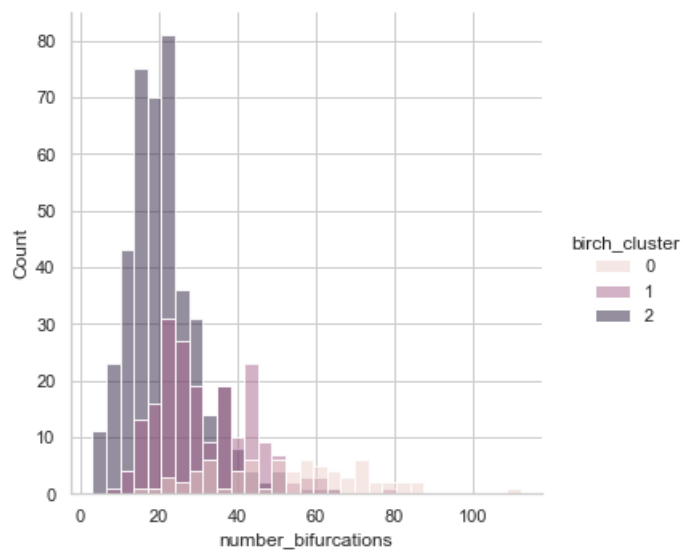


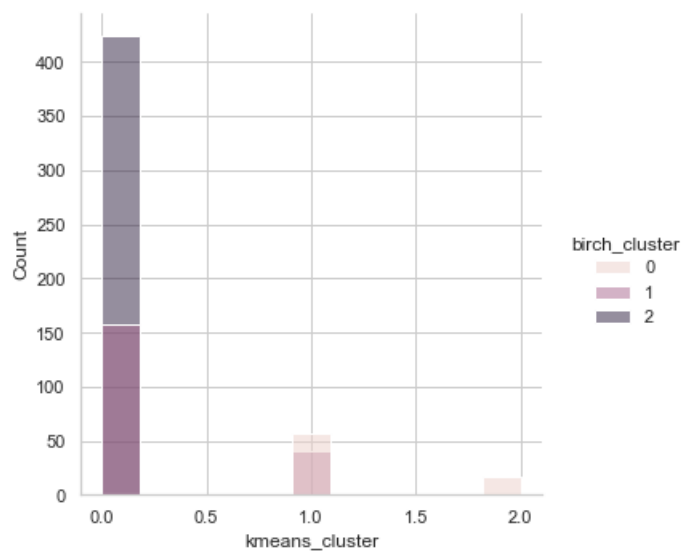
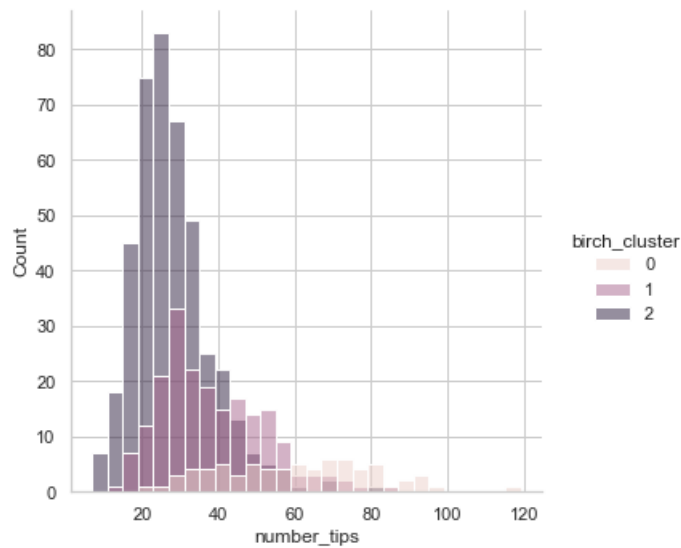
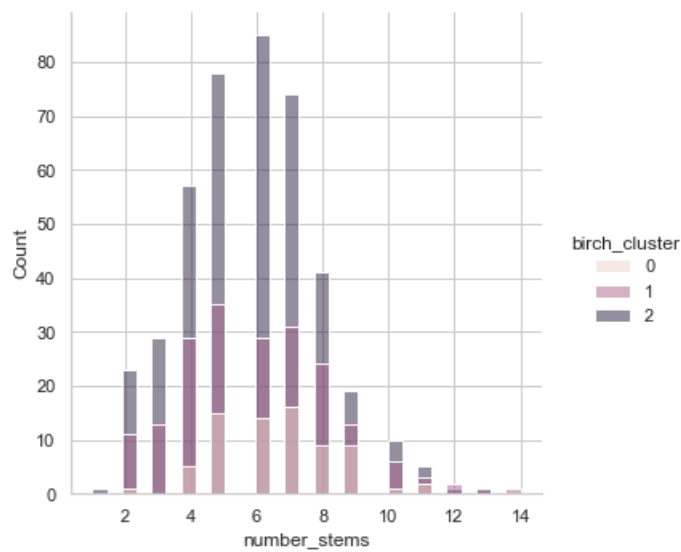












From the scatterplot below, we can see clearly that the discriminatory variables (total length, total surface, total volume, and number nodes) are mostly difference between clusters (not overlapping between each other).

We obtained the same result as what we have assessed in K Clustering analysis part.

```
In [52]: sns.set(style="whitegrid")

fig,axs = plt.subplots(3,2, figsize = (18,18))

#Scatterplot between number nodes and total length measurements
sns.scatterplot(data = df_cluster,
                x = "total_length",
                y = "total_surface",
```

```
hue = 'birch_cluster',
ax=axes[0, 0])

#Scatterplot between number tips and number branches measurements
sns.scatterplot(data=df_cluster,
x="total_length",
y="total_volume",
hue = 'birch_cluster',
ax=axes[0, 1])

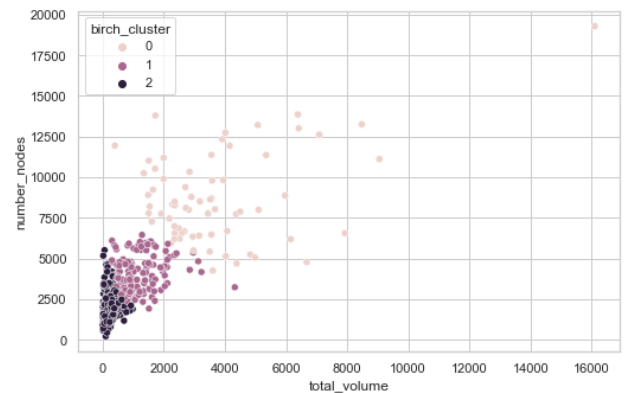
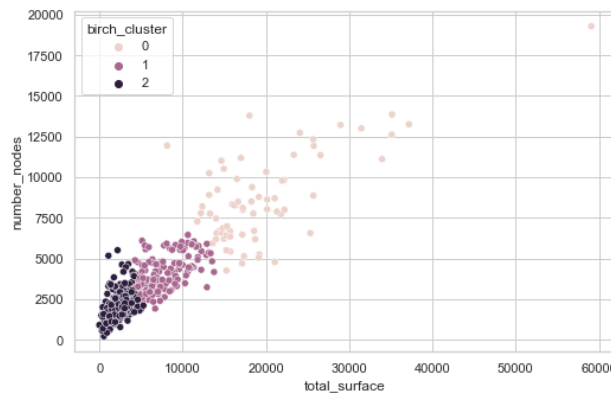
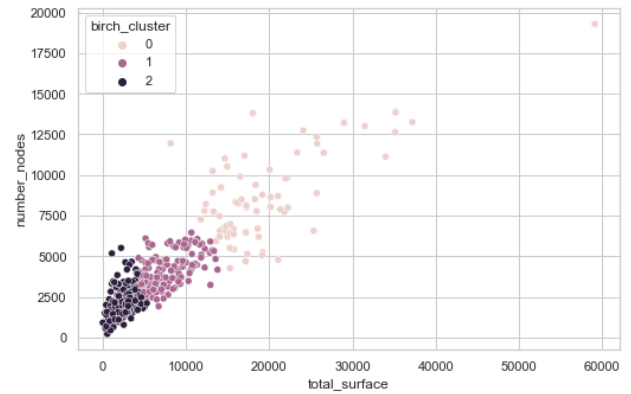
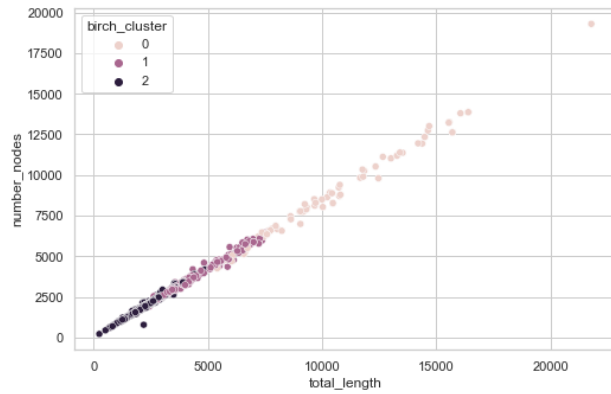
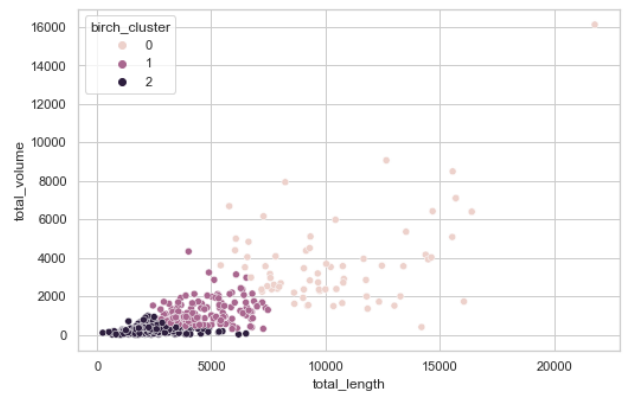
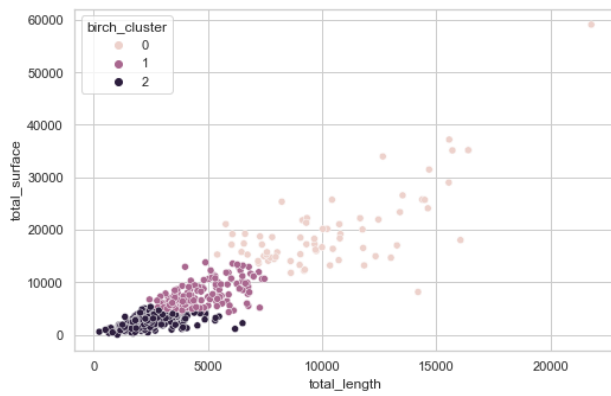
#Scatterplot between number branches and number bifurcations measurements
sns.scatterplot(data=df_cluster,
x="total_length",
y="number_nodes",
hue = 'birch_cluster',
ax=axes[1, 0])

#Scatterplot between number tips and number bifurcations measurements
sns.scatterplot(data=df_cluster,
x="total_surface",
y="number_nodes",
hue = 'birch_cluster',
ax=axes[1, 1])

#Scatterplot between number branches and number bifurcations measurements
sns.scatterplot(data=df_cluster,
x="total_surface",
y="number_nodes",
hue = 'birch_cluster',
ax=axes[2, 0])

#Scatterplot between number tips and number bifurcations measurements
sns.scatterplot(data=df_cluster,
x="total_volume",
y="number_nodes",
hue = 'birch_cluster',
ax=axes[2, 1])

plt.show()
```

References

- ¹ "T-Test: What It Is With Multiple Formulas and When to Use Them." Investopedia, 20 July 2022, www.investopedia.com/terms/t/t-test.asp.
- ² "AIC/BIC in Model Selection." AIC/BIC in Model Selection, www.linkedin.com/pulse/aicbic-model-selection-richard-randa. Accessed 10 Dec. 2022.
- ³ Bevans, Rebecca. "Akaike Information Criterion | When and How to Use It (Example)." Scribbr, 26 Mar. 2020, www.scribbr.com/statistics/akaike-information-criterion.
- ⁴ "Why Does Mean Squared Error Decrease When the Number of Trees Is Increased in Random Forest?" Stack Overflow, 10 May 2019, stackoverflow.com/questions/56084637/why-does-mean-squared-error-decrease-when-the-number-of-trees-is-increased-in-ra
- ⁵ Pramoditha, Rukshan. "Why Do We Set a Random State in Machine Learning Models?" Medium, 2 May 2022, towardsdatascience.com/why-do-we-set-a-random-state-in-machine-learning-models-bb2dc68d8431.
- ⁶ Durukan, Emre. "K-Means Clustering in Python." Medium, 9 July 2022, medium.com/swlh/k-means-clustering-in-python-6c2d7ea01af1.
- ⁷ Maklin, Cory. "BIRCH Clustering Algorithm Example in Python." Medium, 9 May 2022, towardsdatascience.com/machine-learning-birch-clustering-algorithm-clearly-explained-fb9838cbeed9.

"I confirm that all work submitted is my own and that I have neither given, sought, nor received aid in relation to this assignment."

CEM